# Forecasting for the Grain the Decision Is Made At: A Retail World Model on M5, Dominick's and Breakfast at the Frat

Karthik Iyer

Correspondence: karthik.iyer488@gmail.com ORCID: 0009-0004-0593-1602

**Abstract.** Which forecasting model is best depends on the level of the product and store hierarchy at which the forecast is judged, and on how much of the sales the plan (price, promotion, events) moves. We compare a world model, which carries each item's state forward one period at a time under the plan, with a direct lift model, LightGBM, a graph model, seasonal naive, and the pretrained TimesFM 2.5 with and without the plan, on three public data sets: M5 (four 28-day windows), six Dominick's categories and Breakfast at the Frat. Every model is scored at every level of the hierarchy, and every model and change that lost is reported. On M5, TimesFM given only each item's sales history is the most accurate model item by store in all four months (0.836 to 0.885), but forecasts 5% to 10% below total sales; the world model is the most accurate at the chain total in two of the four months and LightGBM in one. The official M5 test month is the only one of the four in which LightGBM is the most accurate over all levels, so a single test month would have given the wrong ranking. On the weekly, promotion-driven data, TimesFM without the plan forecasts 29% to 63% of promoted-week sales and is behind both plan-aware trained models item by store in six of seven data sets; given the plan, it forecasts 60% to 113% and becomes the most accurate model in one category. For an item whose first price cut in a year sold 120 units, the world model forecast 121, 600, 243 across three seeds, against about 3 for LightGBM and 8 for TimesFM: plan-aware models respond in the right direction, and the data do not fix the size. A retailer should choose its forecasting model by the decision the forecast feeds, scored at that decision's level and over several periods.

**Keywords:** retail forecasting; hierarchical evaluation; promotions; world models; foundation models; M5.

## 1 Introduction

A retailer uses one sales forecast for several decisions, and those decisions are made at different levels of the product and store hierarchy. Price changes, promotions, supplier funding and store orders are decided for one item in one store, week by week. Sales, unit and margin plans are set and reviewed for a category or a whole chain, by month. The financial plan is a company total. A forecast can be accurate item by store and still miss the month's category total by several percent, and the reverse: small errors in the same direction on many items add up at the top of the hierarchy, while large errors in opposite directions cancel.

Forecasting competitions and most published comparisons rank models on one combined score over one test period. The M5 competition [1] already weights its score across twelve levels of the Walmart hierarchy, but reports one number per entry, for one 28-day month. This paper asks three questions that a single number hides: which model is most accurate at each level, whether that answer holds from one month to the next, and how much it depends on whether the model sees the plan (the prices, promotions, displays and events set in advance for the weeks being forecast).

We answer them with seven models on three public data sets. The models range from seasonal naive to a pretrained time-series foundation model (TimesFM 2.5 [36]) used without training on our data, and include a world model of our own: each item's recent history is turned into a state, and that state is carried forward one period at a time under the plan, so that what is planned for one week can change the forecast for the weeks after it (Figure 1). Before any final test we ran a development study on windows that precede the final weeks; every configuration in it is reported here, with the ones that lost.

The paper finds and contributes the following.

1. **The most accurate model depends on the level.** On M5, TimesFM given only sales history is the most accurate item by store in all four months tested, by 0.005 to 0.014 against the world model and 0.018 to 0.037 against LightGBM, while models that see the plan are more accurate at every level from store by department up to the chain total (§7.1).
2. **One test month is not enough.** The official M5 test month is the only one of four in which LightGBM is the most accurate over all levels. The world model is lowest in March and, by 0.002, in April, and the

graph model on the practice window. Between the world model and LightGBM, the month's winner is whichever gets the month's total closer (§7.2).

3. **The value of the plan depends on how much of the sales it moves.** Where prices rarely change (M5), giving TimesFM the plan changes its item-by-store error by at most 0.002. Where promoted weeks carry 18% to 61% of a category's units (Dominick's), TimesFM without the plan forecasts 29% to 63% of promoted-week sales, and with it 60% to 113% (§7.3).
4. **Plan-aware models respond to price changes that history has not shown, with a size the data do not fix.** For an item's first price cut in a year the world model forecast 121, 600, 243 units across seeds against 120 sold, where models working from history forecast under 10 (§7.4).
5. **A merchant-readable development record.** The development study finds, among other things, that most of a snack-cracker promotion's lift is taken from other crackers while analgesic and frozen-entree promotions mostly add to the category (§6.2); that a pretrained latent state carries most of the forecast signal without fine-tuning (§6.1); and that development-window rankings carried to the final weeks in some places and reversed in others (§7.5).
6. **Every model and change that lost is reported**, with its cause where one was measured (§8, Table 9).

The paper is organised as follows. Section 2 places the work among retail promotion models, hierarchical evaluation, global and pretrained forecasters, and world models. Section 3 describes the data and the levels scored. Section 4 gives the models and Section 5 the protocol. Section 6 reports the development study, Section 7 the final tests, and Section 8 the failures. Sections 9 to 11 discuss what the results mean for a retailer, their limits, and conclude.

## 2 Related work

### 2.1 Retail forecasting and the promotion plan

Fildes, Ma and Kolassa [5] review retail forecasting research and practice and note that most studies forecast and evaluate at one level of aggregation, usually the item, although retail decisions are spread across levels. Sales of a single item in a single store are intermittent: most items sell nothing on most days, which led to dedicated methods for intermittent demand [6, 7] and to evaluating forecasts as count distributions [8]. The Tweedie family [9], a compound Poisson–gamma distribution with a point mass at zero, is the loss many M5 entries used for this reason, and the loss we use.

Promotions move retail sales more than anything else in the plan. PromoCast forecast promoted-week sales for a chain from the planned event and the item's past response [10]. Van Heerde, Leeflang and Wittink [11] and Macé and Neslin [12] measured the dips in sales before and after a promotion from store scanner data, when shoppers stock up and then buy less. Ma, Fildes and Huang [13] showed that promotional information from within and across categories improves item-level forecasts. Where the extra sales of a promoted item come from is a long-standing question: Gupta [14] split the promotion bump into brand switching, purchase timing and quantity, and van Heerde, Gupta and Wittink [15] showed that measured in units, only about a third of the bump is switching from other brands. Attraction (market-share) models make an item's sales the category's sales times its share, with the share moved by price and promotion [16, 17]; our brand-switching step (§4) follows that structure. Dominick's Finer Foods ran pricing experiments across its stores [18], and its scanner data are behind much of this work [19].

### 2.2 Hierarchies and evaluation

Forecasts that must add up across a hierarchy are studied as hierarchical forecasting, from bottom-up and top-down to optimal reconciliation [21, 22, 23]. We do not reconcile: each model forecasts item by store and its forecasts are summed, which is how the M5 accuracy competition scored entries, with a weighted, scaled error averaged over twelve levels [1]. The scaling by each series' in-sample error follows Hyndman and Koehler's mean absolute scaled error [24]. The M5 data are a reasonable sample of retail series [3], and a companion competition scored forecast distributions [2]. Tashman [25] sets out why a comparison should use several forecast origins rather than one test period, which §7.2 confirms on M5.

### 2.3 Global models, trees and neural forecasters

Models fitted across many series at once, global models, are often more accurate than one model per series [26]. In M5 the leading entries were built on gradient-boosted trees [27, 28]; the winner averaged direct and recursive LightGBM models fitted by store, category and department [4]. Our LightGBM follows that design without the ensembles. Neural global models include N-BEATS [29], DeepAR [30] and the Temporal Fusion Transformer [31], the last two taking known future inputs such as price. Forecasts from different models can

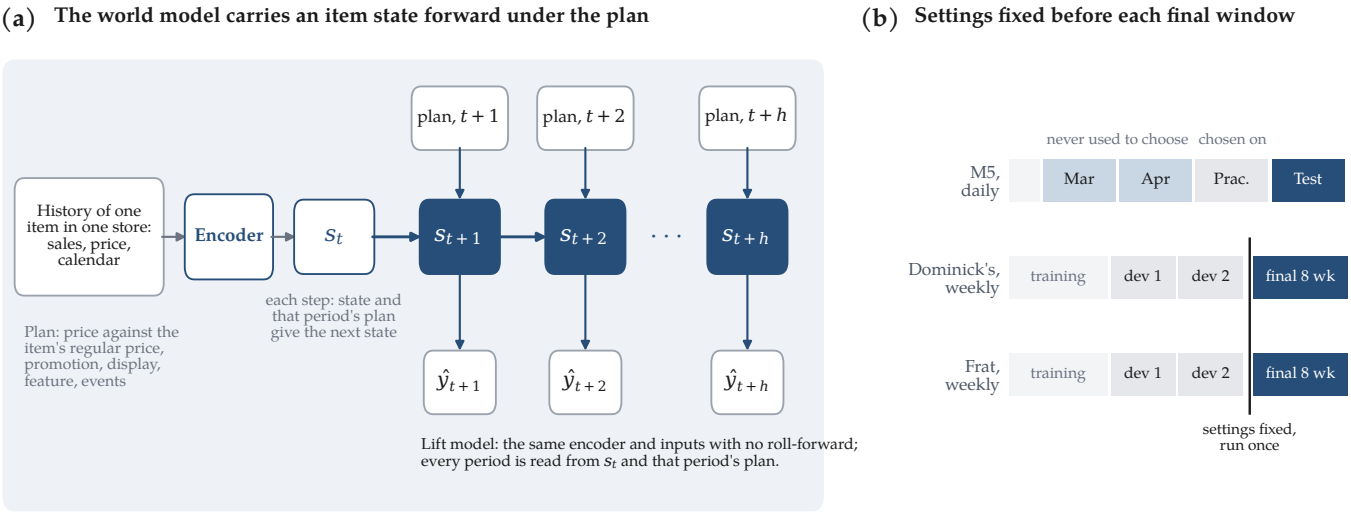



**Figure 1: The world model carries each item's state forward under the plan, and every final window was run once with settings fixed beforehand.** (**a**) An encoder turns the recent history of one item in one store into a state $s_t$. A transition step takes the state and one period's plan and gives the state one period later; each period's sales are read from the state as it stands after the plan so far. The lift model shares the encoder and inputs and reads every period from $s_t$. (**b**) On M5, settings were chosen on the practice window (25 April to 22 May 2016); the two earlier months were not used for any choice. On Dominick's and Frat, settings were chosen on two development windows before the final 8 weeks.

be combined to reduce error [32, 33]; we test combinations in development. Graph neural networks share information between related series; our graph model uses GraphSAGE-style message passing [34], and our item-to-item experiments use attention [59] and conditional similarity networks [35], which learn several similarity spaces between items.

### 2.4 Pretrained foundation models

Pretrained time-series models forecast series they were not trained on from their history alone. TimesFM [36] is a decoder-only transformer pretrained on a large corpus of real and synthetic series; Chronos [38] tokenises values and trains a language model on them; Lag-Llama [39], Moirai [40] and TimeGPT [41] follow similar lines. TimesFM 2.5 adds a covariate method that fits a linear model on future inputs within the batch and forecasts the remainder with the pretrained model [37], which we use to give it the plan. We test on three data sets whether a pretrained model forecasting from history, with or without the plan, matches a model trained on the plan.

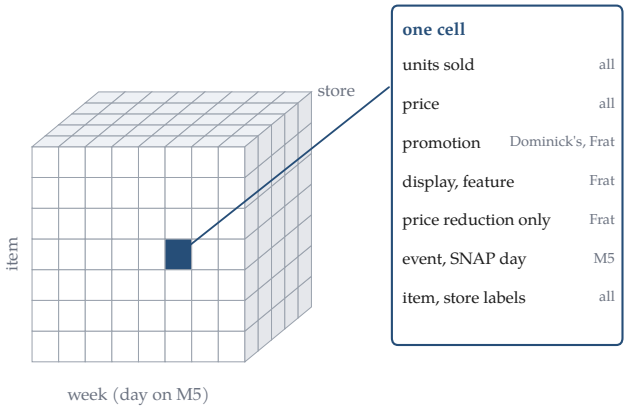
### 2.5 World models and latent prediction

A world model learns a state of its environment and how that state moves under actions, so that the effect of a planned action can be rolled forward before it is taken [42, 43, 44]. LeCun [45] proposed joint-embedding predictive architectures (JEPA), which predict the representation of the future rather than the future itself; I-JEPA [46] and V-JEPA 2 [47] apply the idea to images and video, the latter with an action-conditioned predictor used for planning. Such models need a guard against a collapsed representation: a slowly updated copy of the encoder as the target [48], a variance and covariance penalty [49], or the single isotropic-Gaussian penalty of LeJEPA [50], also used in LeWorldModel [51]. TD-MPC2 [52] rolls a latent state forward under actions and weights near steps more in its loss. DWM [53] keeps the effect of an action apart from the world's own dynamics in the latent state, and AD-WM [54] trains world models to tell actions apart for counterfactual planning. Stable-worldmodel [55] provides a platform for evaluating world models under controlled changes, the same idea as our unusual-plan tests. Closest to us, WorldTS [56] learns covariate-conditioned latent dynamics for forecasting with known future inputs. Our earlier work, M-JEPA [57], compared JEPA fine-tuning recipes one by one with leakage-resistant splits; we carry that method into the development study. In retail the actions are the plan. Our world model applies the structure to sales with a recurrent transition [58] and layer normalisation [60] on top of a transformer encoder [59], and is compared with the same model without the roll-forward so that what carrying the state adds is measured on its own.

## 3 Data and the levels scored

All three data sets are public (Table 1) and share one table shape (Figure 2a). *M5* holds daily Walmart sales

(a) One row per item, store and period




(b) The levels a forecast is scored at, and who plans there

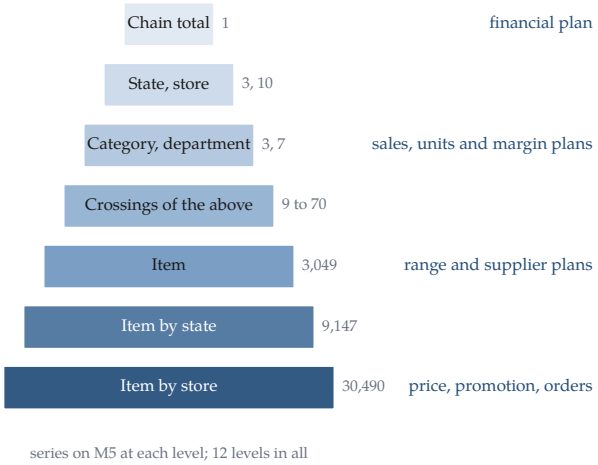



**Figure 2: Every data set is one table of item, store and period, scored at every level a retailer plans at.** (**a**) Each cell holds the units sold, the price and whatever plan the data set records; a lever a data set does not record is passed to the models as unknown, not as zero. (**b**) The M5 hierarchy with the number of series at each level, and the decisions usually made at each. Dominick's and Frat have six levels: total, category, store, category by store, item and item by store.

**Table 1: Three public data sets, all scored item by store and at every level above it.**

| Data | Series | Plan inputs | Final windows |
|---|---|---|---|
| M5 (daily) | 30,490 | Price, event days, SNAP days | Four 28-day windows, 29 Feb to 19 Jun 2016 |
| Dominick's (weekly) | Six categories | Price, promotion | Last 8 weeks of each |
| Frat (weekly) | 3,909 | Price, promotion, display, feature, price reduction only | Last 8 weeks, 16 Nov 2011 to 4 Jan 2012 |

**Table 2: Seven models: five see the plan, two forecast from history alone.** Trained: fitted on the data set; the TimesFM variants are used as published.

| Model | Plan | Trained | What it is |
|---|---|---|---|
| World model | yes | yes | Item state rolled forward under the plan |
| Lift model | yes | yes | Same encoder; baseline times plan lift; no roll-forward |
| LightGBM | yes | yes | One tree model per store, Tweedie loss |
| Graph model | yes | yes | LightGBM's inputs; message passing between related series |
| TimesFM with the plan | yes | no | TimesFM 2.5 with its covariate method |
| TimesFM | no | no | TimesFM 2.5, 512 periods of units |
| Seasonal naive | no | no | Sales one season earlier |

of 3,049 items in 10 stores across three states from 2011 to 2016, with prices by week, event days and the days on which SNAP food assistance is paid [1]. *Dominick's Finer Foods* holds weekly scanner data by store for a Chicago chain from 1989 to 1997, with price, a promotion code and cost [19]. We use six categories chosen before any final run (cheese, snack crackers, soft drinks, bottled juice, analgesics, frozen entrees) and report all 28 categories on the development windows (§6.2). *Breakfast at the Frat* holds weekly sales of 55 breakfast items in 77 stores from 2009 to 2012, with the shelf price, base price, and whether the item was on display, in the weekly circular (feature) or on a tag-only price reduction [20]. No raw data is redistributed.

**Measure.** We use the weighted root mean squared scaled error (WRMSSE) of M5 [1] at each level of the hierarchy: 12 levels on M5 (Figure 2b) and six on Dominick's and Frat. Each series' error is scaled by its own in-sample one-step error, and series are weighted by their dollar sales over the last weeks of training. We report the score over all levels (the M5 score), the chain total and item by store, the ratio of total forecast to total sales, and, on Dominick's and Frat, the ratio of forecast to sales in promoted weeks, in the weeks just after a promotion ends and in ordinary weeks. Lower error is better.

## 4 Models

Seven models are compared in the final tests (Table 2). Five see the plan for the forecast periods and two do not.

**Encoder and state.** For item $i$ in one store, an encoder (a transformer over the last $L$ periods) reads sales and price, each scaled by the item's recent level $\ell_i$, together with the calendar and the recorded levers, and gives a state $s_t$; the item and store labels enter the readout. Sales are scaled so that one model serves items selling one unit a month and items selling hundreds a day. The item's level is its average over the history window, or, on Dominick's, over its last 8 weeks (§6.2).

**World model.** A gated recurrent transition $f$ carries the state forward under the plan $p_{t+h}$ for each future period,

$$s_{t+h}(p) = \mathrm{LN}(f(p_{t+h},\, s_{t+h-1}(p))), \quad s_t(p) = s_t, \qquad (1)$$

where LN is layer normalisation (used on M5 and Dominick's; §6). The plan holds the period's price as a share of the item's regular price (its highest price in the last 12 periods) and the recorded levers. Expected sales are a baseline under the "nothing planned" plan $p^0$ (regular price, no lever) times a lift,

$$\begin{aligned}\log \mu_{i,t+h} = \; & \log \ell_i + b(s_{t+h}(p^0), p^0_{t+h}) \\ & + g(s_{t+h}(p), p_{t+h}) - g(s_{t+h}(p^0), p^0_{t+h}),\end{aligned} \qquad (2)$$

with $b$ and $g$ small networks that also take the item's labels and the horizon step, and the log forecast bounded to $[-10, 10]$. The lift is exactly zero when nothing is planned, and the baseline does not move when the plan changes; both properties are checked by tests in the repository. Training minimises a Tweedie loss on sales relative to $\ell_i$, with batches of item, store and date drawn at random.

**Brand switching (Dominick's).** Items forecast one by one each add their own promotion lift, but in a store and week where several items are promoted much of each item's lift is taken from the others. Following attraction models [16, 17], the world model learns on the last 8 training weeks one number per run, the share $\alpha$ of the summed lift that is new to the category, and in each store and week rescales its item forecasts so that the category total is the baseline total plus $\alpha$ times the summed lift. Each item keeps its share of the total as forecast. On Dominick's three fits with different seeds are averaged in each run. Table 3 gives the settings of the final runs.

**Lift model.** The same model without the roll-forward: $s_{t+h}(p) = s_t$ for every $h$. The two share inputs and training, so the difference between them is the roll-forward. On Frat its readout uses dropout [61] and decoupled weight decay [62], which made it steadier across seeds in development.

**Comparison models.** *LightGBM* [27]: one direct model per store, Tweedie loss, sales at lags no shorter than the horizon, rolling averages and spread of sales, price against the item's highest recent price, calendar and plan inputs, 800 boosting rounds on M5. *Graph model*: LightGBM's inputs, with one round of GraphSAGE-style message passing [34] between the same item in other stores and items of the same department in the same store, Tweedie loss, 6,000 steps. *TimesFM* [36, 37]: the 200M-parameter TimesFM 2.5 checkpoint given each series' last 512 periods of units; we use its mean output, because its default median forecasts 27% to 31% below M5 totals (§8). *TimesFM with the plan*: the same model through its covariate method, a linear model on the plan inputs fitted over groups of 2,048 series, then TimesFM on what that model leaves. *Seasonal naive*: sales one season earlier (a week on M5).

**Table 3: The world model as run in the final tests.** Read from the run manifests in `results/promoted/`. Periods are days on M5 and weeks on Dominick's and Frat.

| Setting | M5 | Dominick's | Frat |
|---|---|---|---|
| Periods of history read | 112 | 26 | 26 |
| Periods of training data | 730 | 156 | 90 |
| State size | 128 | 64 | 64 |
| Encoder layers | 3 | 2 | 2 |
| Attention heads | 4 | 4 | 4 |
| Training steps | 30000 | 6000 | 3000 |
| Batch size | 1024 | 512 | 256 |
| Learning rate | 0.001 | 0.001 | 0.001 |
| Normalised roll-forward | yes | yes | no |
| Level from the last 8 weeks | no | yes | no |
| Brand switching | no | yes | no |
| Fits averaged per run | 1 | 3 | 1 |

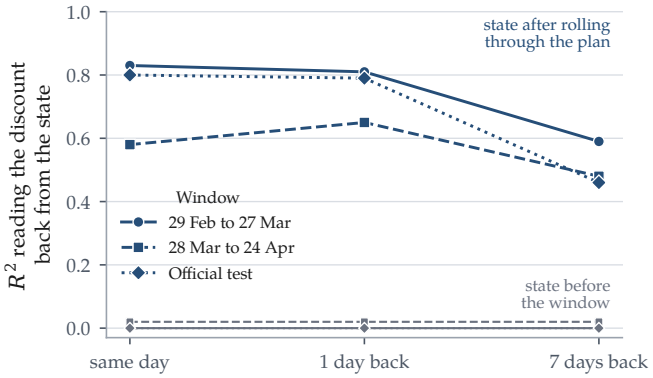



**Figure 3: The rolled state holds the plan, and still holds most of it a week later.** $R^2$ of a ridge read-out of the day's discount (price against regular price) from the frozen world model's state, fitted on half of 6,000 item-store pairs and scored on the other half. Blue: the state after rolling through the window's plan. Grey: the state before the window, which has not seen the plan. Seed 1 in three M5 windows.

**What the state holds.** A probe on the frozen world model shows that the carried state takes in the plan (Figure 3). The day's discount can be read back from the rolled state with $R^2$ 0.58 to 0.83 across three M5 windows, against 0.00 to 0.02 from the state before the window; event and SNAP days at 0.98 to 0.99; and the discount still at 0.46 to 0.59 seven days later. This shows that the plan is in the state; it does not show that the forecast uses it well, which the results test.

## 5 Protocol

**Development and final windows.** On Dominick's and Frat, the last 8 weeks of each data set are the final window, touched once. Development used the two 8-week windows before it, each forecast from everything earlier (Figure 1b). On M5, development used the practice window, 25 April to 22 May 2016, the 28 days before the official test month. After the final tests, two earlier M5 months, 29 February to 27 March and 28 March to 24 April 2016, were run with settings unchanged; they were not used for any choice. One exception is recorded: on Frat the three benchmarks of §6.1 were first scored on three windows that include the final 8 weeks, before that window was locked; none was tuned afterwards.

**Rules fixed in advance.** A component of our model was kept only if removing it made the development score worse; a change that doubles the training cost was kept only if it lowered the score by more than the seed-to-seed range. A benchmark received every input our model received. Each final test ran three seeds per trained model, and the range is reported. The final tests and their settings were written down on 9 October 2026 before any of them ran, as copies of the development configurations with only the name and the scored window changed, which a test in the repository checks. Each was run once and is reported as it came out.

**Models added after the final tests.** TimesFM, TimesFM with the plan and the graph model were added after the final tests as comparisons, with settings fixed before their first run and no development windows of their own.

## 6 Development study

This section reports what was tried before the final tests, on the development windows only. It shows how the final models were chosen and which ideas from the literature did not help on these data. Every row is a section of the results record, with its run folders.

### 6.1 Breakfast at the Frat

**Benchmarks.** On the three windows before the data set's end, seasonal naive scores 0.832, the average of the last 8 weeks 0.756, LightGBM with price only 0.554 and LightGBM with every recorded lever 0.447. The levers are worth a fifth of the error: display, feature and price reduction carry information price does not.

**What helped and what did not.** Figure 4 shows every configuration screened on Frat. Our item encoder with a direct readout scored 0.424, level with LightGBM (0.421). Building the forecast as a baseline under the nothing-planned plan times a lift (the lift model) brought it to 0.406, better than LightGBM in all five runs and with half the spread across seeds. Letting the 55 items in a store affect each other, by conditional similarity, attention or both, made the forecast worse in every version (0.443 to 0.476) and less steady; limiting it to items in the same category narrowed the gap without closing it. Extra inputs did not help: the plan of the four weeks before (0.411), the same week last year (0.408), the US seasonal calendar (0.425), the count of items on promotion in the store (0.406) and the recorded base price used as the regular price (0.434). The recorded base price changes about twelve times a year per item and equals the shelf price in every week without a promotion, so it is not a slow-moving regular price.

**Latent (JEPA-style) losses.** Following the JEPA literature [45, 46] and M-JEPA [57], we added a loss that predicts the encoder's own state for the forecast weeks from the state now and the plan, with a slowly updated copy of the encoder as the target [48]. No recipe beat the lift model: alongside the forecast loss 0.410 to 0.411, pretrained on the latent loss and then fine-tuned 0.407 to 0.414. With the encoder frozen after latent-only pretraining the model still reached 0.428, so a state learned without any sales target carries most of the forecast signal. Long pretraining followed by a slowed encoder gave the best display-week error of any configuration (0.813 for the world model, against 0.817 for LightGBM).

**The world model.** Rolling the state forward week by week under the plan scored 0.408, level with the lift model on the overall score and on every lever. Three ideas from the world-model literature, each at one untuned setting, did not help: keeping the plan's effect apart from the no-plan state as in DWM [53] (0.417), a latent check on each rolled step with a VICReg penalty [49] and near weeks weighted more as in TD-MPC2 [52] (0.427), and all three together (0.424). Training changes did not help either: stopping at the best point on the last 8 training weeks (0.425) took 8 of only 90 training weeks out of the targets, and a term

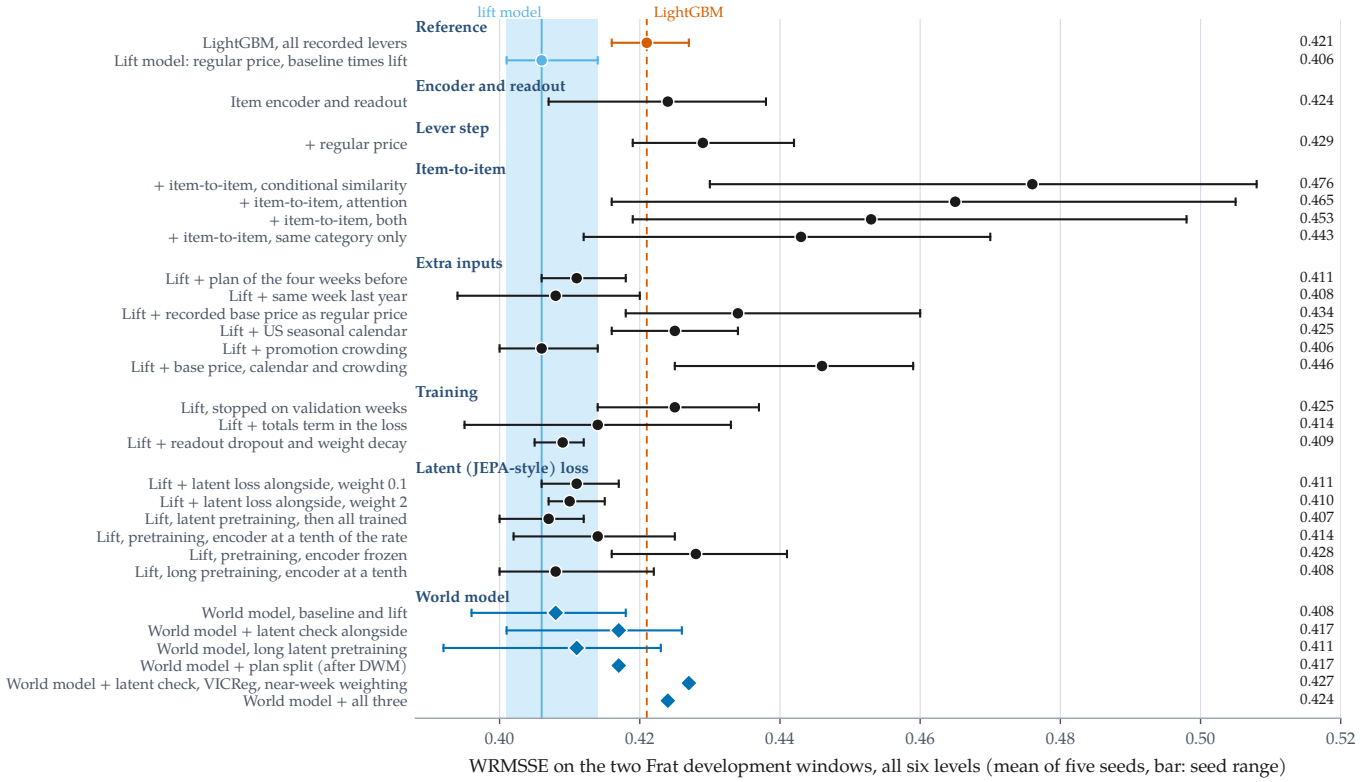



**Figure 4: On Frat the gain came from building the forecast as a baseline times a plan lift; item-to-item links, extra inputs, latent losses and the roll-forward did not improve on it.** Each row is one configuration on the two development windows, five seeds: mean (marker) and range (bar). Blue band: the lift model's seed range; vertical lines: the lift model (0.406) and LightGBM (0.421). Values below the lift model are in bold. Rows without a bar were recorded as a mean only.

on summed totals in the loss (0.414) had no bias to fix, since total forecasts were within 3% of sales in every run. Readout dropout with weight decay (0.409) was level and steadier, and was kept for the lift model.

**Beyond accuracy.** Table 4 reports four further tests. *New items*: with one product in ten held out of training, the lift model borrowing the state and level of the ten known items in the same store with the most labels in common scored 0.890 on the new items, against 1.310 for LightGBM fitted on labels, price and plan, and 1.510 for a forecast of zero. On Dominick's analgesics, where items carry no category or maker labels, borrowing failed (2.999 against 0.923): the method needs item attributes. *Combination*: weighting the lift model and LightGBM on a validation window gave the best Frat result of the study (0.388), better than either part on display and circular weeks, because the two make different errors in promoted weeks; on analgesics the combination was level. Combinations were not carried to the final tests, which compare single models. *Plan A against plan B*: for pairs of weeks of the same item and store with different plans, every model picked the week that sold more about 78% of the time; LightGBM was marginally ahead. *Unusual plans*: on lever mixes an item had run in under 2% of its weeks, no model degraded much and LightGBM stayed ahead. The world model's ability to roll a plan forward did not show as an advantage on these tests.

### 6.2 Dominick's

**Analgesics, the category used to set the models.** On analgesics, seasonal naive scored 1.274, the average of the last 8 weeks 0.723 and LightGBM 0.715. The lift model scored 0.676 on three seeds, and the world model 0.667 once a totals term was removed from its loss: that term made the total worse, pulling forecasts 9% low (0.819 for the lift model). Training four times longer without a stop broke the summed levels (1.473) while item-by-store error barely moved, so a small bias in the same direction on every item added up at the category total. Five seeds changed the picture: the world model held (0.666) and LightGBM scored 0.702, but the lift model fell to 0.718, so a three-seed figure had flattered it.

**All 28 categories.** Run on all 28 categories with settings from analgesics, LightGBM was the most accurate of the three in 14, the lift model in 5 and the world model in 9, with mean WRMSSE 0.688, 0.784 and 0.747; on promoted weeks LightGBM was most accurate in 24 of 27

**Table 4: What the plan-aware models were also tested on during development: new items, combination with LightGBM, choosing between two plans, and plans an item had rarely run.** Development windows, three seeds (Frat combination: five). WRMSSE, lower is better, unless stated. Plan A against plan B: share of week pairs where the forecast picks the week that sold more, and the average miss on the size of the change (log scale).

| Data | Model | Result |
|---|---|---|
| *New items: one product in ten never seen in training (item by store)* | | |
| Frat | LightGBM, new item from labels, price and plan | 1.310 |
| Frat | Lift model, borrowing from similar items | **0.890** |
| Frat | Forecast of zero | 1.510 |
| Analgesics | LightGBM, from labels, price and plan | **0.923** |
| Analgesics | Lift model, borrowing (no attribute labels) | 2.999 |
| *Combination with LightGBM, weights set on a validation window (all levels)* | | |
| Frat | Lift model + LightGBM | **0.388** |
| Frat | World model + LightGBM | 0.390 |
| Frat | LightGBM; lift model; world model | 0.421; 0.406; 0.408 |
| Analgesics | Lift + LightGBM; world + LightGBM | 0.670; 0.668 |
| *Plan A against plan B: same item and store, two weeks with different plans (Frat)* | | |
| Frat | LightGBM | 78.2% right; miss 0.448 |
| Frat | Lift model | 77.4% right; miss 0.460 |
| Frat | World model | 77.5% right; miss 0.459 |
| Frat | World model + plan split | 77.5% right; miss 0.458 |
| *Unusual plans: a lever mix the item had in under 2% of its weeks (item by store)* | | |
| Frat | LightGBM | 0.696 unusual; 0.697 usual |
| Frat | Lift model | 0.743 unusual; 0.725 usual |
| Frat | World model | 0.730 unusual; 0.730 usual |

(Figure 5, Table 5). Analgesics did not generalise, which is why the paper reports every category. The failures (snack crackers 1.463, bath soap 1.465, the lift model on soft drinks 1.406) sat at the category total: item-by-store error was close to LightGBM's. On a six-store sample of snack crackers the cause was in-and-out and seasonal lines: items recorded in 8 to 20 of the last 26 weeks were 8% of sales and the lift model forecast them at 9 times their sales, because the model set each item's level from its 26-week average after the item had fallen away. Capping extreme weeks did little; taking the level from the last 8 weeks removed most of it. On all 28 categories the 8-week level brought the world model's mean to 0.717, removed the large failures (snack crackers 1.463 to 0.874, bath soap 1.465 to 1.195) and gave a little back in categories that had been fine, so it beat LightGBM in 10 categories against 14 before. LightGBM, which sees short recent averages directly, does not have the problem.

**Peaks and dips.** In cheese, both our model and LightGBM under-forecast promoted weeks by about a third in the holiday window, and cheese peaks do not sit on a fixed week relative to the holidays: they follow the chain's promotion timing, which event flags cannot place. Figure 6 splits forecast against sales into promoted weeks, the weeks just after a promotion, and ordinary weeks, for six categories. The world model did not see the dip after a promotion where the dip is large: it forecast 1.33 times sales after a promotion in snack crackers and 1.32 in soft drinks, against 0.92 and 0.90 for LightGBM. Feeding the model's own sales forecast into the roll-forward, so that a promotion's sales could shape the weeks after it, fixed the soft-drink dip (1.32 to 1.00) and made the overall score worse in four of six categories. Weighting the loss by dollar sales and by how far a week sits from the item's usual level fixed cheese promoted weeks (1.21 to 1.00) and was worse overall in three categories. In frozen entrees every model, LightGBM included, forecast about 0.7 of sales after a promotion: there the week after is stronger than any model expected.

**From the first world model to the one tested.** Figure 7 follows the world model on the six categories through the changes kept, three seeds each. Across seeds the first version was unsteady: its range within a category was 0.04 to 0.38, against under 0.03 for LightGBM, and on the mean LightGBM was ahead in five of six. Brand switching fixed what it was built for: in snack crackers and soft drinks, where promoted weeks broke the total, it was better on every seed (0.813 to 0.752, 0.902 to 0.789) and item-by-store error improved too. The share of the

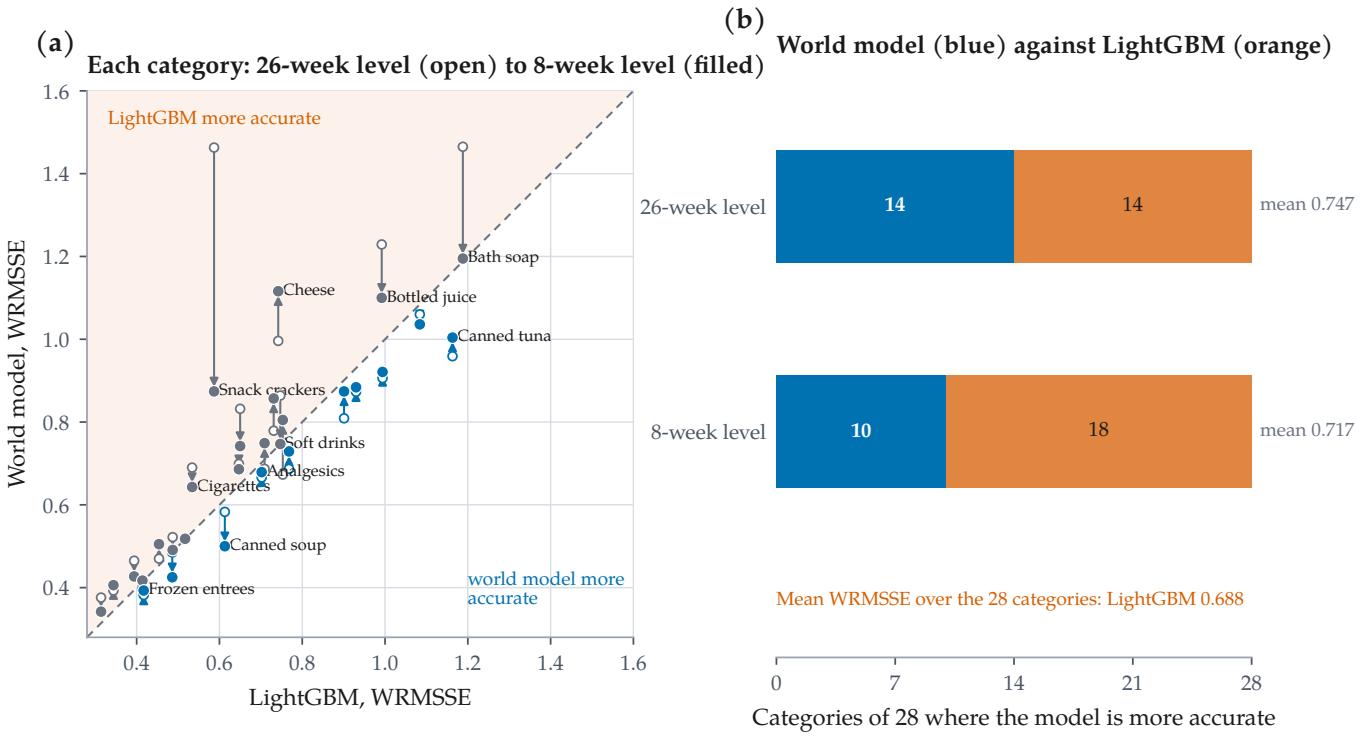



**Figure 5: Run unchanged on all 28 Dominick's categories, the world model set on analgesics failed badly in a few; taking each item's level from its last 8 weeks removed the worst failures but did not make it better than LightGBM on average.** Development windows, one seed. **(a)** Each category's WRMSSE for LightGBM (horizontal) and the world model (vertical), with the item level from the last 26 weeks (open) and the last 8 weeks (filled); blue: the world model is more accurate with the 8-week level. **(b)** Categories where each model is more accurate, and the mean over the 28.

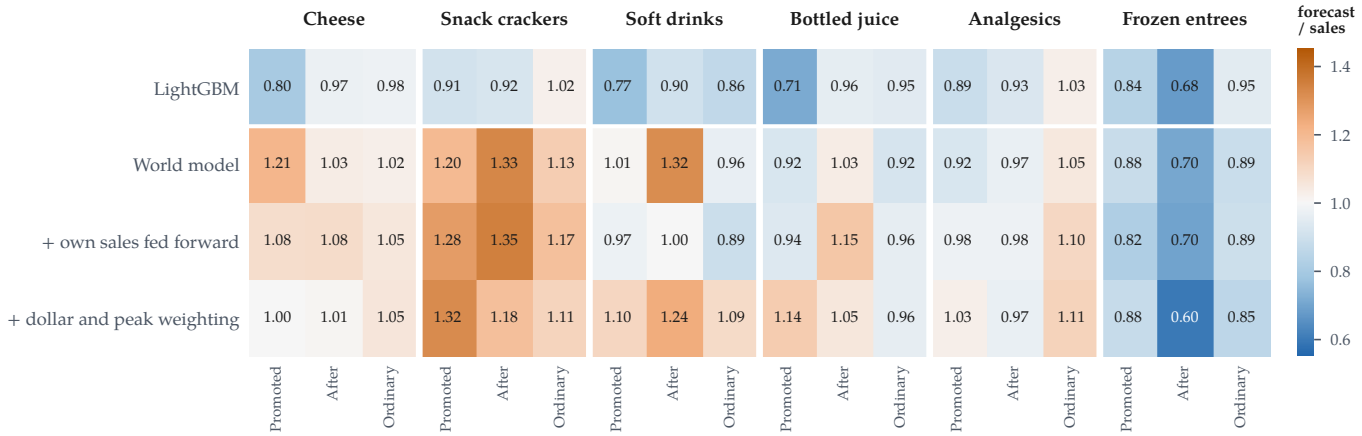



**Figure 6: The world model over-forecast the weeks after a promotion where shoppers had stocked up; the two fixes each repaired one category and hurt others.** Forecast against sales (1.0 is right; orange over, blue under) in promoted weeks, the weeks just after a promotion and ordinary weeks, Dominick's development windows, one seed. Rows: LightGBM, the world model with the 8-week level, with its own sales forecast fed forward, and with the loss weighted by dollar sales and by peaks.

summed lift that is new to the category, learned per run, reads as a merchant would expect (Figure 7g): in snack crackers 0.00 to 0.66, so most of a promotion's lift is taken from other crackers, against about 0.9 in bottled juice, analgesics and frozen entrees. Where it comes out at 1.0 the step changes nothing, a built-in check. Van Heerde, Gupta and Wittink [15] found a similar split from household data; here it is learned from store sales alone. Normalising the rolled state narrowed the seed range in five of six categories, and averaging three fits per run with normalisation gave the steadiest version (range 0.05 or less in five categories) and the candidate for the final test: better than LightGBM on the three-seed mean in analgesics (0.622 against 0.701), frozen entrees (0.391 against 0.419) and soft drinks (0.743 against 0.759). One snack-cracker run of the three failed (1.806): its second win-

**Table 5: Across all 28 Dominick's categories LightGBM is more accurate on average; the world model wins in 14 with its first level and 10 with the 8-week level.** Development windows, one seed (analgesics: five), settings set on analgesics and not tuned per category. World model with the item's level from its last 26 weeks or last 8 weeks. Lowest in each row in bold.

| Category | LightGBM | Lift | World, 26 wk | World, 8 wk |
|---|---|---|---|---|
| Analgesics | 0.702 | 0.718 | **0.666** | 0.679 |
| Bath soap | **1.188** | 1.568 | 1.465 | 1.195 |
| Beer | **0.454** | 0.460 | 0.470 | 0.505 |
| Bottled juice | **0.992** | 1.248 | 1.229 | 1.100 |
| Canned soup | 0.613 | 0.532 | 0.583 | **0.500** |
| Canned tuna | 1.163 | 1.030 | **0.959** | 1.004 |
| Cereals | 0.753 | 0.795 | **0.673** | 0.805 |
| Cheese | **0.742** | 0.900 | 0.996 | 1.116 |
| Cigarettes | **0.534** | 0.918 | 0.690 | 0.643 |
| Cookies | **0.487** | 0.503 | 0.522 | 0.491 |
| Crackers | 0.709 | 0.967 | **0.687** | 0.749 |
| Dish detergent | **0.344** | 0.436 | 0.393 | 0.406 |
| Fabric softener | 0.414 | **0.358** | 0.397 | 0.417 |
| Front-end candies | **0.650** | 0.748 | 0.832 | 0.742 |
| Frozen dinners | 0.930 | 0.905 | **0.873** | 0.884 |
| Frozen entrees | 0.417 | **0.374** | 0.382 | 0.393 |
| Frozen juices | **0.314** | 0.360 | 0.376 | 0.342 |
| Grooming products | 0.901 | 1.208 | **0.809** | 0.874 |
| Laundry detergent | 0.768 | 0.735 | **0.687** | 0.729 |
| Oatmeal | 1.084 | **1.026** | 1.060 | 1.036 |
| Paper towels | 0.486 | 0.445 | 0.485 | **0.425** |
| Shampoos | 0.994 | 0.938 | **0.906** | 0.921 |
| Snack crackers | **0.587** | 0.959 | 1.463 | 0.874 |
| Soaps | 0.517 | 0.521 | **0.516** | 0.518 |
| Soft drinks | **0.747** | 1.406 | 0.864 | **0.747** |
| Toilet tissue | **0.394** | 0.413 | 0.465 | 0.427 |
| Toothbrushes | **0.647** | 0.671 | 0.699 | 0.686 |
| Toothpaste | **0.731** | 0.802 | 0.779 | 0.857 |
| Mean | 0.688 | 0.784 | 0.747 | 0.717 |

dow's category total broke with promoted weeks forecast at 2.6 times sales while item-by-store error stayed normal, so the summed promotion lift can still run away on one fit in three. Twice the training steps, screened after the final tests were fixed, was better in two categories, worse in two, level in one and mixed in one, and was not kept.

**Other changes.** Showing the model both its 26-week and its recent levels as inputs, instead of setting its scale by the recent level, was the best world-model version in four of ten categories but did not fix snack crackers (1.515) or soft drinks (1.069): the model did not learn to switch to the recent level when the long level still set its scale. Item role labels (new line, in-and-out, volume driver, long tail, high-low, promotion-responsive, core), assigned from each item's own sales before the windows, moved LightGBM by at most 0.013, since it reads the same information from its history, and moved the world model in both directions; letting the role pick the level broke snack crackers (2.422). We did not call the volume drivers key value items, since a key-value-item list rests on shopper price perception, which these data do not record.

**Table 6: On the M5 practice window, only the normalised roll-forward separated the world model from LightGBM; no later change beat it.** WRMSSE over all 12 levels, 25 April to 22 May 2016. A screened change was kept only if it lowered the score, and a change that doubles the cost only if by more than the seed range (0.029).

| Model or change | Mean | Seed range |
|---|---|---|
| Seasonal naive | 0.870 | |
| LightGBM, one model per store | 0.715 | |
| Lift model (3 seeds) | 0.713 | 0.683 to 0.734 |
| Item encoder and readout (5 seeds) | 0.699 | 0.669 to 0.725 |
| World model (3 seeds) | 0.707 | 0.680 to 0.750 |
| World model + normalised roll-forward (3 seeds) | 0.656 | 0.646 to 0.675 |
| + average of three fits (1 run) | 0.665 | |
| + weekday as seven inputs (seed 1) | 0.653 | |
| + weekday and event kinds (seed 1) | 0.646 | |
| + 60,000 training steps (seed 1) | 0.642 | |
| + item level against past-year level (seed 1) | 0.680 | |

### 6.3 M5

On the M5 practice window (Table 6), the item encoder with a direct readout scored 0.699 over five seeds against 0.715 for LightGBM, and the lift and world models were level with LightGBM over three seeds (0.713 and 0.707). Normalising the rolled state was the one change that separated the world model: 0.656, with its worst seed (0.675) better than the best seed without it (0.680), all of the gain at the totals and the seed range cut from 0.070 to 0.029. It helped on M5's 28 daily steps and not on Dominick's 8 weekly steps, which fits the reading that it stops the state drifting over a long roll. Averaging three fits did not help on M5 (0.665), where single fits were already steady.

After the official test month, an error analysis of the world model found its 2.6% excess in Foods (+3.9%), in one store (CA_3, +7.5%), and on Sundays and Mondays (+5.1%, +6.4%), not growing with the horizon. Grouped by sales in the 28 days before the test, the excess came from items that had been selling fast (+5.4% and +6.1% on items selling 3 to 10 and over 10 a day). Four changes were screened on the practice window under rules written before screening: weekday as seven inputs, event kinds, twice the training steps, and the item's level against its past-year level. None beat the original (Table 6), so there was no second version. On the prac-

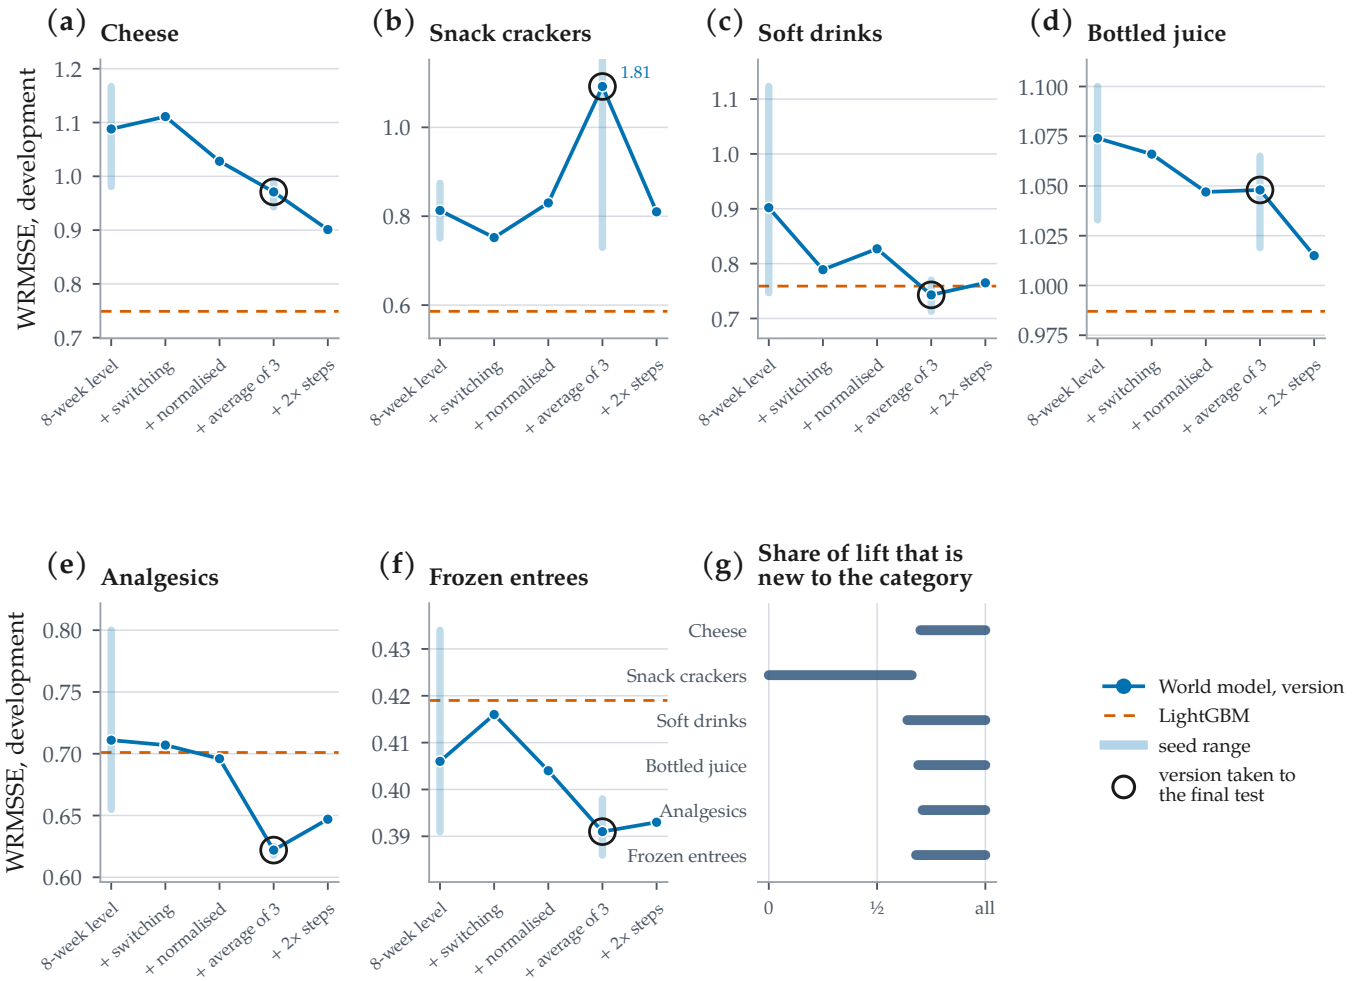



**Figure 7: Brand switching, a normalised roll-forward and averaging three fits took the world model past LightGBM in three of six categories on the development windows, and the switching step learned how much of a promotion's lift is new to the category.** **(a–f)** Development WRMSSE of each version, three-seed means; shaded bars: seed range of the first version and of the one taken to the final test (circled); dashed: LightGBM. **(g)** The learned share of the summed promotion lift that is new to the category (the rest is switching between items), range over runs.

tice window every version under-forecast by 4% to 6%, while on the test month the world model over-forecast by 2.6%: the sign of the level error changed between adjacent months, so a fix tuned on one would not carry to the next. That finding is what led to the confirmation months of §7.2.

## 7 Final results

The results follow the order of the contributions: the level at which a model is judged (§7.1), the number of months it is judged on (§7.2), the plan (§7.3, §7.4), and what carried over from development (§7.5). Table 7 and Table 8 hold every final score.

### 7.1 The most accurate model depends on the level

On M5, TimesFM given only each item's sales history is the most accurate model item by store in all four months (Figure 8): 0.836, 0.828, 0.828 and 0.885, against 0.851, 0.833, 0.834 and 0.897 for the world model and 0.873, 0.846, 0.847 and 0.905 for LightGBM. Among the trained models, the world model and the lift model are ahead of LightGBM and the graph model item by store in every month, and in both confirmation months every world-model seed is ahead of every LightGBM seed.

At the total the order reverses. TimesFM forecasts 5% to 10% below total sales in every month, so its total error is higher than seasonal naive's in three of the four months (0.830 against 0.658 in March) and level with it on the test month (0.558 against 0.560). The world model has the lowest total error in March (0.398) and April (0.283), the graph model on the practice window (0.455) and LightGBM on the test month (0.247).

**Table 7: On M5 the best model depends on the level scored and on the month.** WRMSSE on the four 28-day windows (lower is better), mean over seeds; seeds per window in the second column. All: the 12 levels together; Total: the chain total; I×S: item by store. The lowest value in each column is in bold. TimesFM with the plan has no recorded total. The lift model's April mean includes one failed seed (2.603 overall; §8).

| | | 29 Feb to 27 Mar | | | 28 Mar to 24 Apr | | | 25 Apr to 22 May (practice) | | | 23 May to 19 Jun (test) | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| **Model** | **Seeds** | All | Total | I×S | All | Total | I×S | All | Total | I×S | All | Total | I×S |
| World model | 3/3/3/3 | **0.642** | **0.398** | 0.851 | **0.555** | **0.283** | 0.833 | 0.656 | 0.535 | 0.834 | 0.638 | 0.379 | 0.897 |
| Lift model | 3/3/3/3 | 0.664 | 0.427 | 0.846 | 1.246 | 0.964 | 0.834 | 0.713 | 0.623 | 0.832 | 0.663 | 0.368 | 0.896 |
| LightGBM | 3/3/1/3 | 0.709 | 0.578 | 0.873 | 0.557 | 0.351 | 0.846 | 0.715 | 0.659 | 0.847 | **0.547** | **0.247** | 0.905 |
| Graph model | 3/3/3/3 | 0.778 | 0.621 | 0.878 | 0.644 | 0.493 | 0.845 | **0.641** | **0.455** | 0.851 | 0.673 | 0.412 | 0.909 |
| TimesFM | 1/1/1/1 | 0.854 | 0.830 | **0.836** | 0.819 | 0.786 | **0.828** | 0.896 | 0.930 | **0.828** | 0.729 | 0.558 | 0.885 |
| TimesFM with the plan | 1/1/1/1 | 0.858 | – | 0.837 | 0.772 | – | 0.829 | 0.929 | – | 0.830 | 0.719 | – | **0.883** |
| Seasonal naive | 1/1/1/1 | 0.889 | 0.658 | 1.121 | 0.923 | 0.668 | 1.126 | 0.870 | 0.751 | 1.100 | 0.847 | 0.560 | 1.176 |

**Table 8: On the promotion-driven weekly data, models that see the plan lead, and which one leads changes by category.** WRMSSE on the final 8 weeks (lower is better). World model, lift model and LightGBM: mean of three seeds, settings fixed on the development windows. TimesFM and TimesFM with the plan: one run each, added after the final tests. The lift model was run on Frat only. The lowest value in each row and group is in bold.

| | All levels | | | | | Item by store | | | | |
|---|---|---|---|---|---|---|---|---|---|---|
| **Data** | World | Lift | LightGBM | TimesFM | TimesFM+plan | World | Lift | LightGBM | TimesFM | TimesFM+plan |
| Frat | 0.681 | 0.714 | **0.549** | 0.965 | 0.686 | 0.665 | 0.677 | **0.608** | 0.827 | 0.641 |
| *Dominick's* | | | | | | | | | | |
| Cheese | 0.897 | – | 0.708 | 0.829 | **0.545** | 0.715 | – | 0.702 | 0.792 | **0.699** |
| Snack crackers | 0.448 | – | **0.412** | 0.468 | 0.475 | 0.608 | – | **0.596** | 0.640 | 0.614 |
| Soft drinks | **0.830** | – | 0.989 | 1.070 | 0.974 | 0.734 | – | **0.652** | 0.895 | 0.780 |
| Bottled juice | 0.987 | – | 0.994 | **0.940** | 0.990 | 0.816 | – | **0.750** | 0.782 | 0.754 |
| Analgesics | **0.867** | – | 0.978 | 1.009 | 0.889 | **0.713** | – | 0.714 | 0.790 | 0.742 |
| Frozen entrees | 0.485 | – | **0.478** | 0.563 | 0.566 | 0.557 | – | **0.529** | 0.682 | 0.599 |

Figure 9 shows where the order changes. From the chain total down to store by department a plan-aware model is the most accurate in every month, with the world model's error typically 20% to 50% below TimesFM's. At the three item levels (item, item by state, item by store) the models are within a few percent of each other, and TimesFM is the most accurate in three of the four months, item by store in all four. A model that is best for store orders on M5 is therefore not the model to use for a category plan. The combined M5 score weights the twelve levels equally, and nine of them are above the item, so a ranking on it mostly reflects the upper levels.

### 7.2 One test month misleads

Over all twelve levels, the official M5 test month (23 May to 19 June 2016) is the only one of the four in which LightGBM is the most accurate model: 0.547 against 0.638 for the world model, level with the 10th-ranked entry of the competition (0.547), with the winner at 0.520 and the best statistical benchmark at 0.671. In March the world model is ahead, 0.642 (seeds 0.632 to 0.648) against 0.709 (0.706 to 0.713); in April the two are level, 0.555 against 0.557; and on the practice window the graph model is lowest (0.641), with the world model at 0.656 and LightGBM at 0.715.

The ranking follows the level of the total forecast (Figure 10a): across models, months and seeds, the size of the total bias and the total error have a correlation of 0.87. Between the world model and LightGBM, LightGBM forecasts 3% to 6% below total sales in March and April and 2% below on the test month; the world model is within 1% in March and April and 2.6% above on the test month. In each of these three months the one of the two whose total is closer to sales has the lower total error and the lower error over all levels. Grouping April's item-store pairs by their March sales rate (Figure 10b), both models are within 6% of sales in every group that sold in March; LightGBM is under in every group and the world model is closer on 50.6% to 54.2% of pairs in each. The world model's weak spot is pairs that sold nothing in March, where it forecasts 62% above April

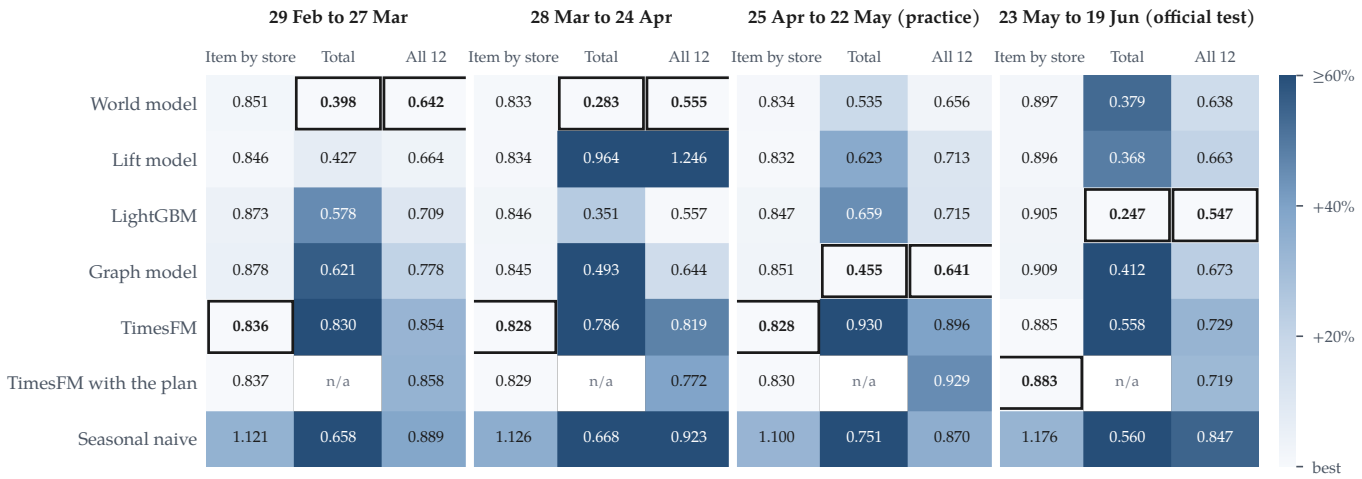



**Figure 8: On M5, TimesFM from history alone has the lowest item-by-store error in every month; at the total and over all levels the plan-aware models do, and which one changes by month.** WRMSSE of every model, four windows, three grains; cell shade: how far above the lowest value in its column (white: the lowest, outlined). Trained models: mean of three seeds (LightGBM on the practice window: one run). The lift model's April mean includes one failed seed (§8). TimesFM with the plan has no recorded total.

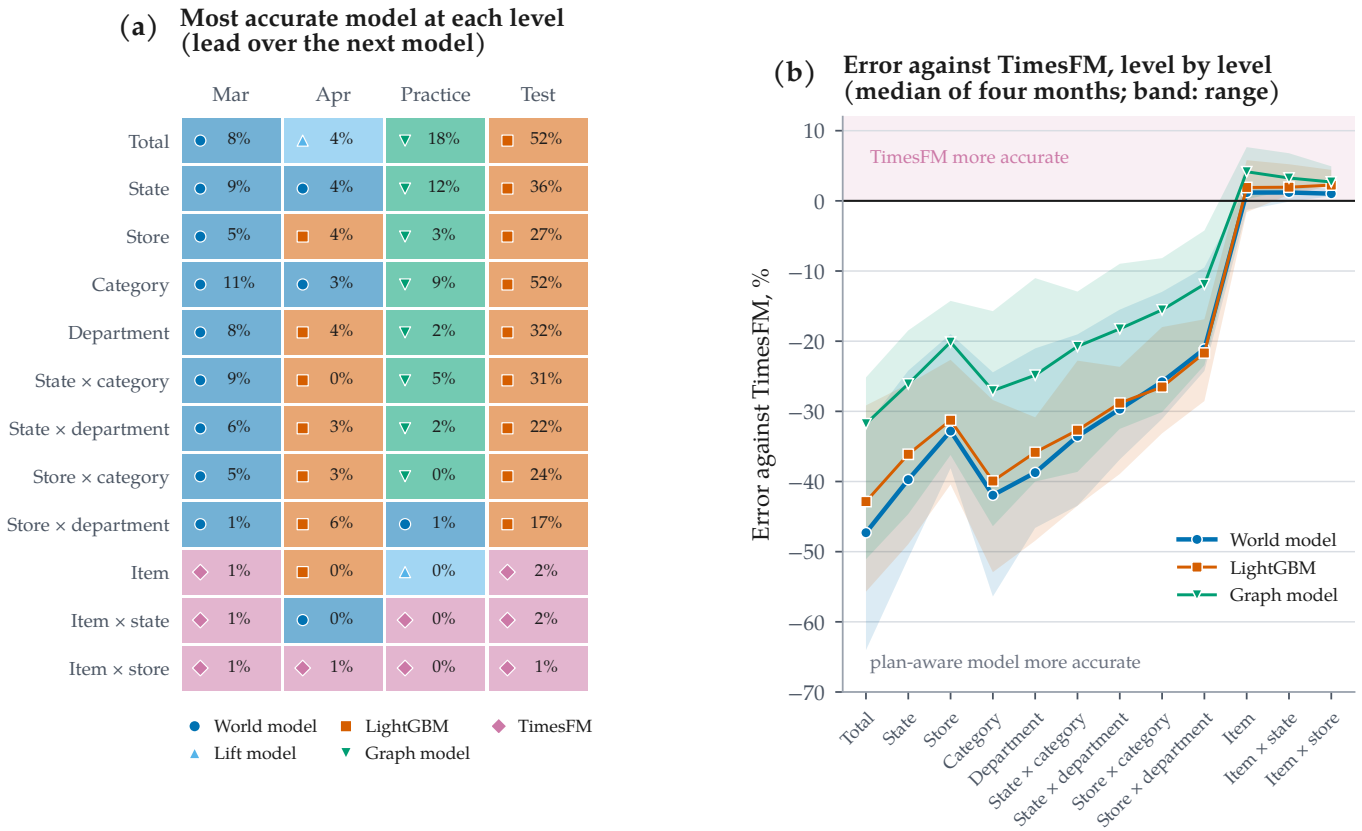



**Figure 9: From the chain total down to store by department a plan-aware model is most accurate; at the three item levels the models are within a few percent and TimesFM usually leads.** **(a)** The most accurate of five models at each M5 level in each window (cell colour and marker), with its lead over the next model in percent (lift model: median of its seeds). **(b)** Error of the world model, LightGBM and the graph model against TimesFM at each level: median of the four windows (line) and range (band); below zero the plan-aware model is more accurate.

sales.

A comparison on the official test month alone would have ranked LightGBM first. On three of the four months it is not.

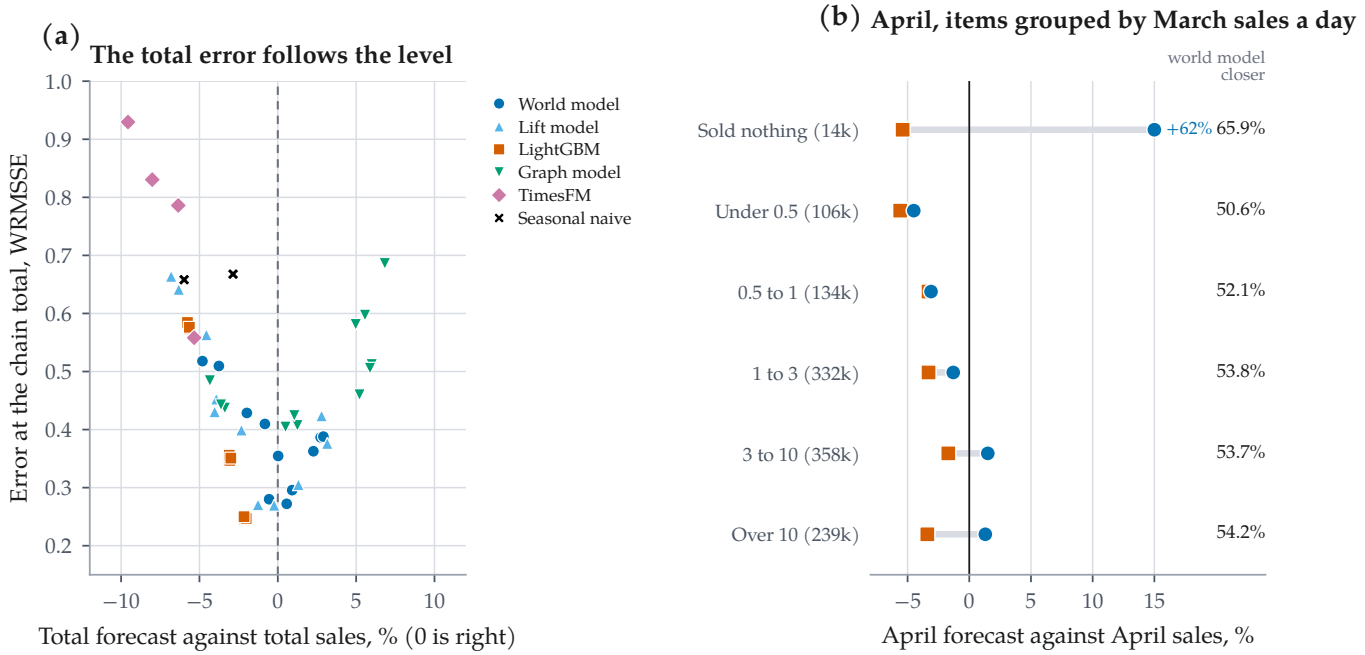



**Figure 10: Whichever model gets the month's total closer has the lower total error, and that changes by month.** (**a**) Error at the chain total against the total forecast's bias, one point per model, M5 month and seed ($n = 50$; one failed lift seed lies off the chart). The correlation of the size of the bias with the error is 0.87. Seasonal naive's bias was not recorded on the two later windows. (**b**) April, with each item-store pair grouped by its average daily sales in March (units sold in April in brackets), so that the grouping does not depend on the sales being scored: forecast against sales for the world model (circle) and LightGBM (square), and the share of pairs where the world model is closer.

### 7.3 How much the plan matters

The plan matters in proportion to how much of the sales it moves. On M5, where prices change rarely, giving TimesFM the plan changes its item-by-store error by at most 0.002 in any month (0.837 against 0.836 in March). On Dominick's and Frat, where promotions drive the weeks, it changes a great deal (Figure 11).

Without the plan, TimesFM is behind both the world model and LightGBM item by store on Frat and in five of six Dominick's categories; in bottled juice it is between them. This is the reverse of M5. The reason shows in the promoted weeks (Figure 12a): TimesFM forecasts 29% to 63% of what sold in a promoted week, because nothing in its input says the promotion is coming.

Given the plan, TimesFM forecasts 60% to 113% of promoted-week sales, and its overall error improves on Frat (0.965 to 0.686) and in three of six categories (cheese, soft drinks, analgesics). It becomes the most accurate model in cheese (0.545 against 0.708 for LightGBM) and is level with the world model on Frat (0.686 against 0.681), but stays behind a model trained on the plan in the other categories. In bottled juice the plan made it slightly worse (0.940 to 0.990).

The trained plan-aware models also fall short in promoted weeks: the world model forecasts 48% to 88% of promoted-week sales and LightGBM 50% to 84%. Among the trained models, which one is more accurate depends on the category (Table 8): over all levels the world model is ahead in soft drinks (0.830 against 0.989) and analgesics (0.867 against 0.978), with every world-model seed ahead of every LightGBM seed in both; level in bottled juice; slightly behind in frozen entrees (0.485 against 0.478); and behind in cheese, snack crackers and on Frat. Item by store LightGBM is level or ahead in every category. On Dominick's the world model's gains are at the totals, the opposite pattern to M5.

### 7.4 First-time price cuts

In the April M5 window, HOUSEHOLD_1_146 had sold almost nothing for a year at 5.98, then sold 120 units in four stores over a window that held a one-week cut to 4.50 (Figure 12b). Models working from history forecast almost nothing: LightGBM 3.1, 2.9, 3.3 units across seeds, TimesFM 7.8. The world model, which reads the cut through its plan, forecast 121, 600, 243; the graph model 59, 45, 55; the lift model 79, 206,886, 316.

Models that read the plan respond to a price change that history has not shown, and LightGBM, which sees price only against its own recent highest, responds little. How large the response should be is not fixed by the data: the world-model seeds differ by a factor of five, and one lift-model seed forecast over 200,000 units. For

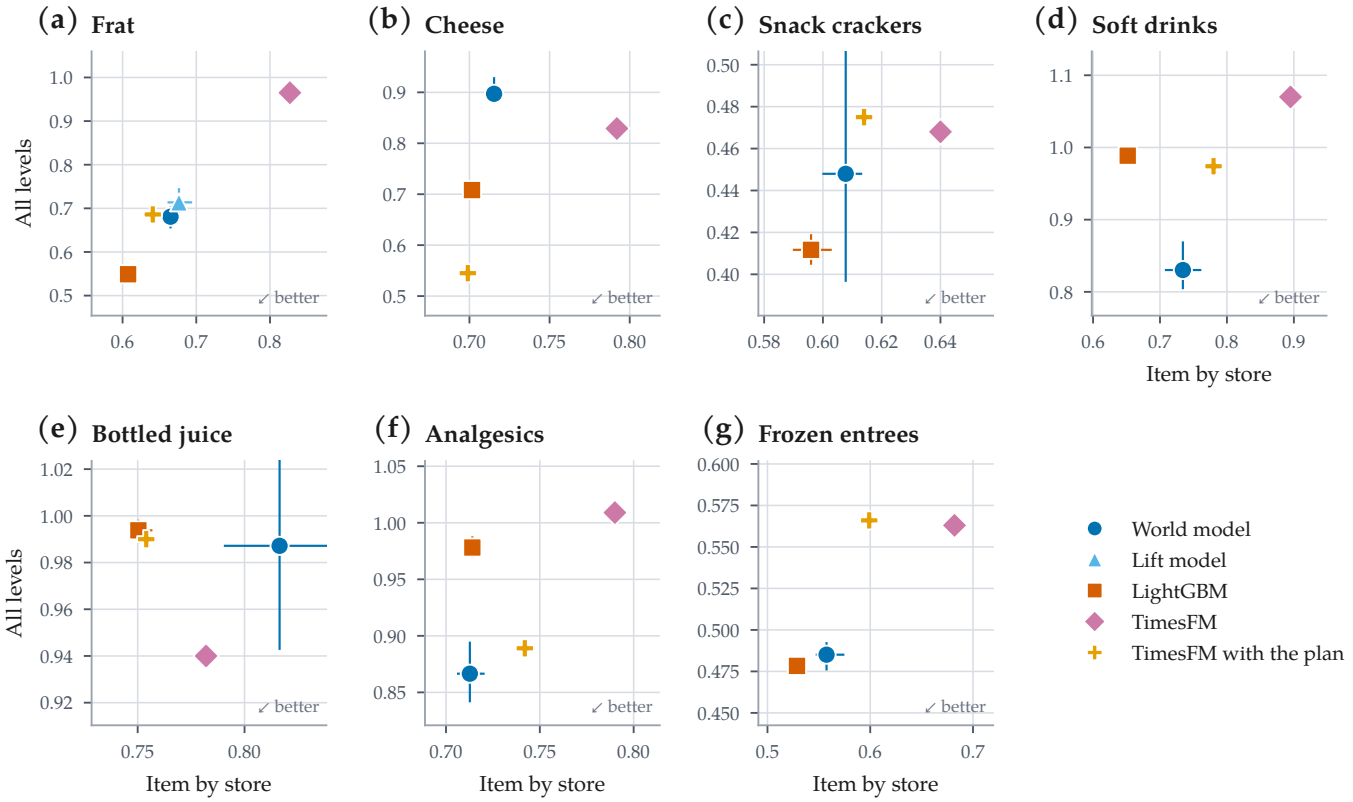



**Figure 11: On the weekly data each data set has its own trade-off between item-by-store and all-level accuracy, and TimesFM without the plan is behind both trained models item by store in six of seven.** Final 8 weeks: item-by-store WRMSSE (horizontal) against all-level WRMSSE (vertical), lower-left is better. World model, LightGBM and (Frat only) lift model: mean and seed range in both directions; TimesFM variants: one run. Axes differ by panel.

the same window, FOODS_3_443 in one store had a 68% cut and sold 4 units; every plan-reading model forecast more (the world model up to 26).

### 7.5 What carried over from development

Development rankings carried to the final weeks in some places and not in others. On Dominick's the world model was ahead of LightGBM in development in analgesics, frozen entrees and soft drinks; in the final weeks it was ahead in analgesics and soft drinks, by more than the seed ranges, and slightly behind in frozen entrees. On M5 the practice-window lead (0.656 against 0.715) held in March, was a tie in April and reversed on the test month. On Frat the three trained models were level in development (0.406 to 0.421) and LightGBM was clearly ahead in the final weeks (§8). In each case the difference between development and final followed what changed in the data: the month's level on M5, the holiday weeks on Frat. A development win is evidence about the windows it was measured on, and the paper reports both.

## 8 Negative and degenerate results

We report the failures at the same length as the results. Two had a measured cause; the other changes and comparison models are in Table 9.

**A lift-model seed that broke on five series.** In April, lift-model seed 2 scored 2.603 over all levels against 0.566 and 0.570 for seeds 1 and 3, and forecast 1,413,724 units against 1,183,626 sold. Five item-store series hold all of the excess. Each had its first price cut in a year inside the window (the four HOUSEHOLD_1_146 stores of §7.4, and FOODS_3_443 in one store, 2.50 to 0.79 for three weeks). Seed 2 read the cuts as lifts up to the model's bound, about 22,000 times the item's level per day. Item by store, weighted by dollar sales, the failure is invisible (0.834); at the totals it dominates. A price planner looking at items would see a few absurd forecasts; a category planner would see the month's total 19% too high. The world model, with the same readout and the roll-forward, did not fail on these series in any seed.

**Frat holiday weeks.** On the Frat development windows the world model, lift model and LightGBM were level (0.406 to 0.421). On the final 8 weeks the world model

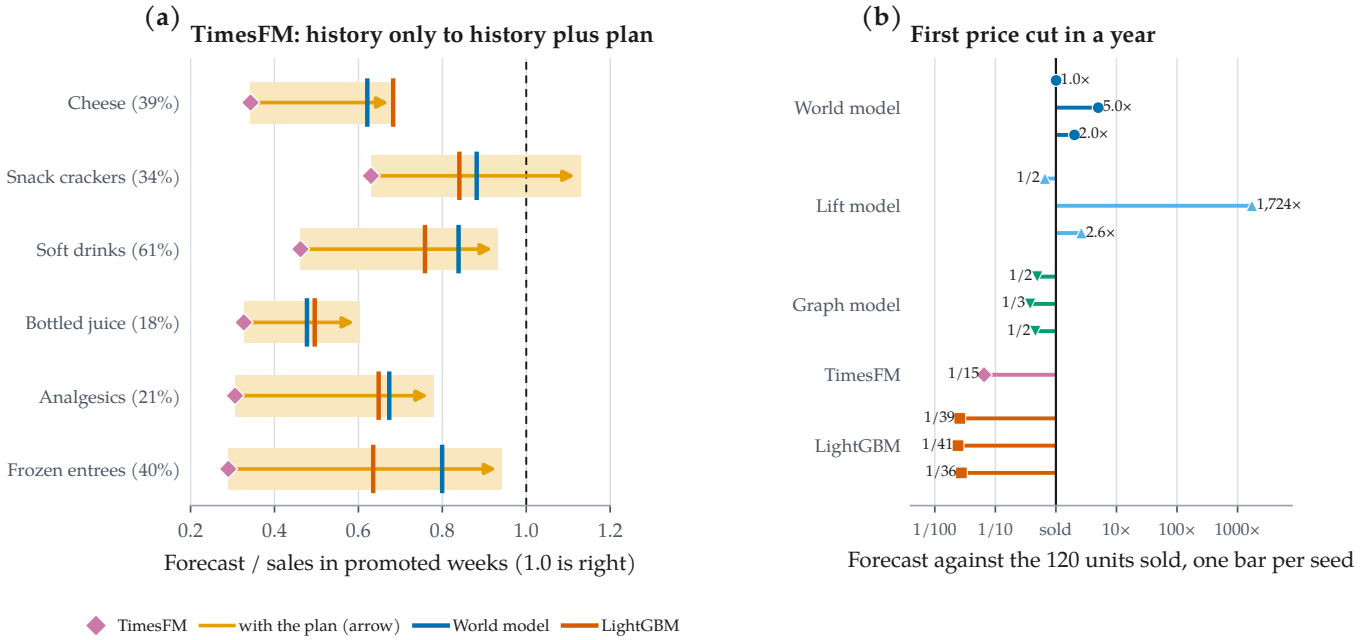



**Figure 12: Given the plan, TimesFM closes most of its promoted-week shortfall; for a price cut history has not shown, plan-aware models respond and history-only models do not.** **(a)** Forecast against sales in the promoted weeks of the final 8 weeks, by Dominick's category (promoted weeks' share of the category's units in brackets): TimesFM from history alone (diamond) to TimesFM with the plan (arrow head), with the world model and LightGBM (ticks). **(b)** HOUSEHOLD_1_146 in the four M5 stores where it had its first price cut in a year (5.98 to 4.50 for one week) inside the April window: each seed's forecast over the 28 days against the 120 units sold, log scale.

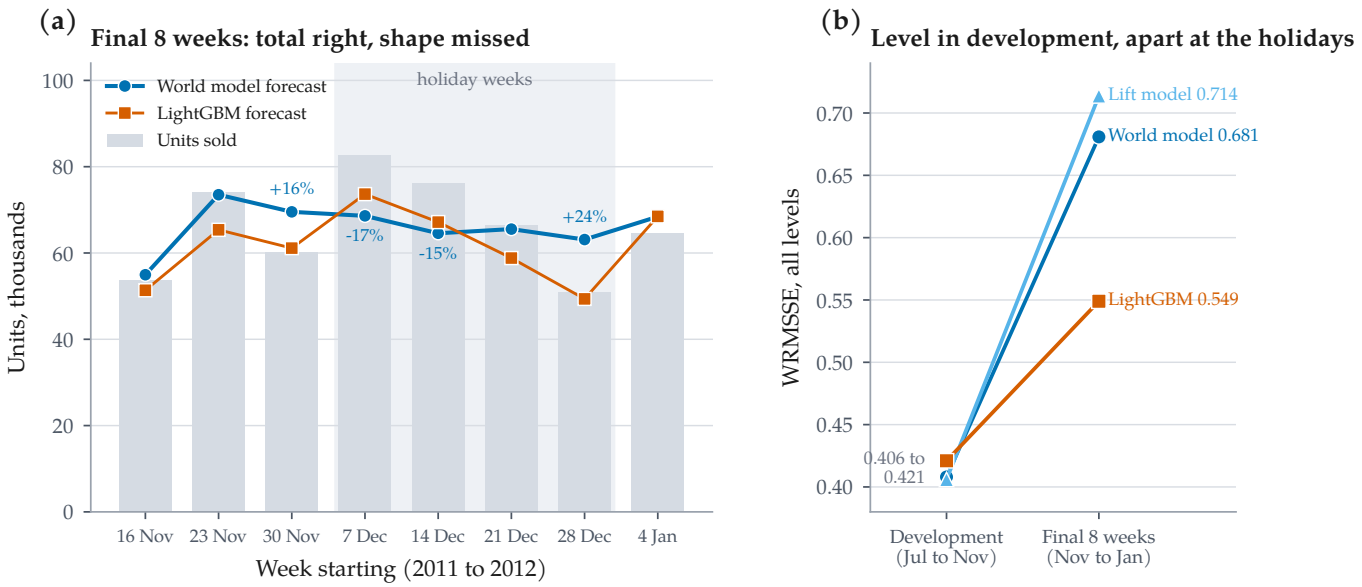



**Figure 13: On Frat the world model gets the 8-week total right and misses the holiday shape, which the development windows did not contain.** **(a)** Units sold each week of the final window over all item-stores (bars), and the weekly forecasts of the world model and LightGBM, three seeds averaged; labels: the world model's largest weekly misses. **(b)** All-level WRMSSE on the development windows and on the final 8 weeks.

fell to 0.681 against 0.549 for LightGBM, and every LightGBM seed beats every world-model seed. The final weeks are the holiday period, which the development windows (late July to mid November) did not contain (Figure 13). The world model's 8-week total is right (forecast against sales 1.000, LightGBM 0.937), but it is 15% to 17% under in the two December peak weeks and 24% over in the week after Christmas. LightGBM, with week-of-year and month inputs and training data that include the previous holiday season, follows that shape more closely. The world model is still ahead of the lift model here (0.681 against 0.714).

**TimesFM's default output.** TimesFM's default point forecast is the median of each day's sales. For an item that sells nothing on most days the median is near zero while the mean is not, so the median forecast is 27% to 31% below M5 totals in every month, and its upper-level error is worse than seasonal naive's (overall 1.71 to 1.88). Item by store it was close to LightGBM. A median is the wrong target for a merchant's units and dollars at category level; the mean output, used throughout, does not have the problem.

**A withdrawn claim.** An earlier reading that the world model was more accurate than LightGBM on fast-selling items came from grouping pairs by their sales in the scored window, which pulls forecasts in high groups towards under and in low groups towards over for both models. Grouped by an earlier window (Figure 10b), the advantage is small and the same in every group. We do not make the claim.

**Changes that did not become defaults.** Table 9 lists every change and comparison model that lost, and §6 gives the development detail. Longer training, weekday and event inputs, a long-run level, a totals term in the loss, a count likelihood, a separate plan path in the state, feeding the model's own forecast into the roll-forward and item role labels either made no difference beyond seed noise or made the model worse. The graph model is level with LightGBM item by store, whose inputs it shares, so sharing information between related series added nothing on these inputs; neither did item-to-item links in our own model on Frat or Dominick's.

## 9 Discussion

A retailer should choose a forecasting model by the decision it feeds, scored at that decision's level and over several periods (Table 10).

**Why the level changes the answer.** Item-by-store error is dominated by the noise of single items selling a few units, which history read well can match and the plan adds little to when prices rarely move. The totals are dominated by the month's level: a 2% bias in the same direction on every item is invisible item by store and decides the category total. TimesFM forecasts from history what an item usually sells and has no input that says this month will differ; the plan-aware models are trained across all items and stores and see events, SNAP days and prices, and get the level closer. On weekly promotion-driven data the plan moves the single items too, so the plan-aware models win item by store as well.

**What the roll-forward adds.** Against the lift model, which differs only in the roll-forward, the world model is level or better over all levels in every M5 month and on Frat, it was steadier across seeds once the rolled state was normalised, and it did not fail on the first-time price cuts that broke a lift-model seed. Item by store it is not more accurate than the lift model, and on M5 neither is more accurate than a pretrained model that sees no plan at all. The development tests designed to show what rolling a plan forward buys (plan A against plan B, unusual plans) did not separate it from LightGBM. On these data, the case for the roll-forward is steadiness and a bounded response to unfamiliar plans, not accuracy on familiar ones.

**Two cautions.** A direct lift readout can extrapolate a price cut it has never seen far beyond any sale, and whether it does depends on the seed; a plan-driven forecast needs a bound and a check on series whose plan falls outside their own history. And a forecast that is accurate item by store can be several percent off the category total, which is the number a merchant plans against.

## 10 Limitations and threats to validity

**Windows and seeds.** M5 has four final windows; Dominick's and Frat one each. On the M5 practice window LightGBM was run once. Much of the development study used one seed, and single-seed world-model runs varied by up to 0.07 on analgesics; the decisions that mattered were confirmed on three to five seeds. Three seeds show the spread of training, not of the data.

**Comparison models added later.** TimesFM, TimesFM with the plan and the graph model were added after the final tests, with settings fixed before their first run but without development windows of their own. Each TimesFM variant is one run. TimesFM with the plan fits its linear plan model within groups of 2,048 series, a memory limit, so the fit is shared within each group.

**What lift means.** Merchants promote in weeks they expect to sell well, so part of a measured lift is good timing. The lift is what the forecast attributes to the plan, not proof of what the plan caused. Dominick's in-store pricing experiments [18] are the place to test it against a cause; we have not done so.

**The size of a plan response.** The world model's response to a first-time price cut varied by a factor of five across seeds. The data do not identify it, and a deployed system would need a bound or a prior.

**Table 9: Every change and comparison model that did not become a default.** Overall WRMSSE unless stated; "against" gives the default it was compared with on the same window. dev.: development windows before the final 8 weeks. Practice: the M5 window 25 April to 22 May 2016. Each row is a section of `results/README.md` with its run folders.

| Change or model | Where tested | Result |
|---|---|---|
| Twice the training steps | Dominick's (6) and Frat, dev., 3 seeds | Better in 2, worse in 3, level or mixed in 2 |
| Weekday as seven inputs | M5 practice, seed 1 | 0.653 against 0.646 |
| Event kinds (with weekday) | M5 practice, seed 1 | 0.646 against 0.646 |
| 60,000 training steps | M5 practice, seed 1 | 0.642 against 0.646; gain below seed range 0.029 |
| Item level against past-year level | M5 practice, seed 1 | 0.680 against 0.646 |
| Totals term in the loss | Dominick's analgesics, dev. | 0.819 against 0.683; total 9% under |
| Negative binomial likelihood | Frat, analgesics, dev. | Level |
| Plan split, latent check, near-week weighting | Frat, dev., 5 seeds | 0.417 to 0.427 against 0.408 |
| Own forecast fed into the roll-forward | Dominick's (6), dev. | Worse overall in 4 of 6 |
| Item role labels | Dominick's (6), dev. | No effect on LightGBM; snack crackers 2.422 for the world model |
| Averaging three fits | M5 practice, seed 1 | 0.665 against 0.656 (kept on Dominick's) |
| Graph model over product attributes | M5, four windows, 3 seeds | Level with LightGBM item by store |
| TimesFM, median output | M5, four windows | Overall 1.71 to 1.88; totals 27% to 31% under |
| Lift model (no roll-forward) | M5, four windows, 3 seeds | One seed failed in April (2.603) |

**Table 10: Each decision has its own level, and the results point to a different model for each.**

| Decision | Level | What the results say |
|---|---|---|
| Store orders and inventory for items whose price rarely moves | Item by store, day or week | A pretrained model from history alone is as accurate as or more accurate than plan-aware models (M5: most accurate in all four months) |
| Promotion, price change, supplier funding | Item by week, under a planned change | A model that reads the plan is needed; history alone forecasts 29% to 63% of promoted-week sales and does not respond to a first price cut. The size of a response to a plan unlike history needs a bound |
| Category sales, units and margin plans | Category or chain, week or month | Judge on the level of the total; the world model was within 1% of total M5 sales in both confirmation months, and on Dominick's its gains were at the totals |
| Promotion planning across a category | Category by store-week | Account for switching between items: in snack crackers most of a promotion's lift came from other crackers |
| Choosing a model | Any | Score at least three or four periods; one period reversed the ranking on M5 |

**What the data do not record.** None of the data sets records stock, so a sell-out looks like low demand; or shelf space, or placement within a display or circular. Weather is not known 8 weeks ahead. The data were collected between 1989 and 2016, and their plan inputs differ, so the results compare models within each data set, not across them.

## 11 Conclusion

Which forecasting model is most accurate depends on the level it is judged at, the month, and how much of the sales the plan moves. On M5 a pretrained model reading history alone is the most accurate item by store in every month, while plan-aware models are more accurate at the totals, and the official test month is the only one of four in which LightGBM wins overall. On promotion-driven data a model without the plan forecasts well under half of promoted-week sales in most categories, and adding the plan to a pretrained model recovers much of that without replacing a model trained on the plan. A world model that carries each item's state forward under the plan responds to price changes history has not shown, in the right direction and with an uncertain size, and learned from store sales how much of a promotion's lift is new to a category. Retail forecasts should be chosen and reported by the decision they feed, at that decision's level, over several periods, with the losing models stated.

## 12 Reproducibility and compute

All code, configuration files and tests are in the repository. Every run records its code commit (runs refuse

to start on uncommitted code), its configuration and checksum, the seed, checksums of every data file read, package versions and the device in a manifest; the manifest and scores of the 110 runs read by this manuscript are committed under `results/promoted/`, and its figures, tables and quoted final-test numbers are generated from them by `paper/figs/make_figures.py` and `paper/figs/make_tables.py`. Numbers from runs not in that folder (the development study, TimesFM on Dominick's and Frat, TimesFM with the plan, and the analyses of saved forecasts) are copied into `paper/figs/rwdata.py` from `results/README.md`, which names each run folder. Forecast files and model checkpoints are kept outside the repository, with their checksums recorded in each run's scores. GPU runs used one NVIDIA Tesla T4 (16 GB) on Google Colab; LightGBM ran on the Colab CPU. The world model took about 100 minutes per seed on M5 and 9 minutes per run (three fits) on a Dominick's category; LightGBM 78 minutes per seed on M5; the graph model 7.5 and TimesFM 4 minutes per M5 window.

## 13 Data and code availability

M5 is available from the competition's Kaggle page [1]; Dominick's from the Kilts Center for Marketing, University of Chicago Booth School of Business [19]; Breakfast at the Frat from dunnhumby [20]. Code: https://github.com/Kar488/retail-world-models.

## References


[1] S. Makridakis, E. Spiliotis, and V. Assimakopoulos. M5 accuracy competition: Results, findings, and conclusions. *International Journal of Forecasting*, 38(4):1346–1364, 2022.

[2] S. Makridakis, E. Spiliotis, V. Assimakopoulos, Z. Chen, A. Gaba, I. Tsetlin, and R. L. Winkler. The M5 uncertainty competition: Results, findings and conclusions. *International Journal of Forecasting*, 38(4):1365–1385, 2022.

[3] E. Theodorou, S. Wang, Y. Kang, E. Spiliotis, S. Makridakis, and V. Assimakopoulos. Exploring the representativeness of the M5 competition data. *International Journal of Forecasting*, 38(4):1500–1506, 2022.

[4] Y. In and J.-Y. Jung. Simple averaging of direct and recursive forecasts via partial pooling using machine learning. *International Journal of Forecasting*, 38(4):1386–1399, 2022.

[5] R. Fildes, S. Ma, and S. Kolassa. Retail forecasting: Research and practice. *International Journal of Forecasting*, 38(4):1283–1318, 2022.

[6] J. D. Croston. Forecasting and stock control for intermittent demands. *Operational Research Quarterly*, 23(3):289–303, 1972.

[7] A. A. Syntetos and J. E. Boylan. The accuracy of intermittent demand estimates. *International Journal of Forecasting*, 21(2):303–314, 2005.

[8] S. Kolassa. Evaluating predictive count data distributions in retail sales forecasting. *International Journal of Forecasting*, 32(3):788–803, 2016.

[9] B. Jørgensen. Exponential dispersion models. *Journal of the Royal Statistical Society, Series B*, 49(2):127–145, 1987.

[10] L. G. Cooper, P. Baron, W. Levy, M. Swisher, and P. Gogos. PromoCast: A new forecasting method for promotion planning. *Marketing Science*, 18(3):301–316, 1999.

[11] H. J. van Heerde, P. S. H. Leeflang, and D. R. Wittink. The estimation of pre- and postpromotion dips with store-level scanner data. *Journal of Marketing Research*, 37(3):383–395, 2000.

[12] S. Macé and S. A. Neslin. The determinants of pre- and postpromotion dips in sales of frequently purchased goods. *Journal of Marketing Research*, 41(3):339–350, 2004.

[13] S. Ma, R. Fildes, and T. Huang. Demand forecasting with high dimensional data: The case of SKU retail sales forecasting with intra- and inter-category promotional information. *European Journal of Operational Research*, 249(1):245–257, 2016.

[14] S. Gupta. Impact of sales promotions on when, what, and how much to buy. *Journal of Marketing Research*, 25(4):342–355, 1988.

[15] H. J. van Heerde, S. Gupta, and D. R. Wittink. Is 75% of the sales promotion bump due to brand switching? No, only 33% is. *Journal of Marketing Research*, 40(4):481–491, 2003.

[16] D. E. Bell, R. L. Keeney, and J. D. C. Little. A market share theorem. *Journal of Marketing Research*, 12(2):136–141, 1975.

[17] L. G. Cooper and M. Nakanishi. *Market-Share Analysis: Evaluating Competitive Marketing Effectiveness*. Kluwer Academic Publishers, Boston, 1988.

[18] S. J. Hoch, B.-D. Kim, A. L. Montgomery, and P. E. Rossi. Determinants of store-level price elasticity. *Journal of Marketing Research*, 32(1):17–29, 1995.

[19] James M. Kilts Center for Marketing. Dominick's dataset. University of Chicago Booth School of Business. https://www.chicagobooth.edu/research/kilts/research-data/dominicks.

[20] dunnhumby. Breakfast at the Frat. dunnhumby source files. https://www.dunnhumby.com/source-files/.

[21] G. Athanasopoulos, R. A. Ahmed, and R. J. Hyndman. Hierarchical forecasts for Australian domestic tourism. *International Journal of Forecasting*, 25(1):146–166, 2009.

[22] R. J. Hyndman, R. A. Ahmed, G. Athanasopoulos, and H. L. Shang. Optimal combination forecasts for hierarchical time series. *Computational Statistics & Data Analysis*, 55(9):2579–2589, 2011.

[23] S. L. Wickramasuriya, G. Athanasopoulos, and R. J. Hyndman. Optimal forecast reconciliation for hierarchical and grouped time series through trace minimization. *Journal of the American Statistical Association*, 114(526):804–819, 2019.

[24] R. J. Hyndman and A. B. Koehler. Another look at measures of forecast accuracy. *International Journal of Forecasting*, 22(4):679–688, 2006.

[25] L. J. Tashman. Out-of-sample tests of forecasting accuracy: An analysis and review. *International Journal of Forecasting*, 16(4):437–450, 2000.

[26] P. Montero-Manso and R. J. Hyndman. Principles and algorithms for forecasting groups of time series: Locality and globality. *International Journal of Forecasting*, 37(4):1632–1653, 2021.

[27] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu. LightGBM: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems 30*, 2017.

[28] T. Januschowski, Y. Wang, K. Torkkola, T. Erkkilä, H. Hasson, and J. Gasthaus. Forecasting with trees. *International Journal of Forecasting*, 38(4):1473–1481, 2022.

[29] B. N. Oreshkin, D. Carpov, N. Chapados, and Y. Bengio. N-BEATS: Neural basis expansion analysis for interpretable time series forecasting. *International Conference on Learning Representations (ICLR)*, 2020.

[30] D. Salinas, V. Flunkert, J. Gasthaus, and T. Januschowski. DeepAR: Probabilistic forecasting with autoregressive recurrent networks. *International Journal of Forecasting*, 36(3):1181–1191, 2020.

[31] B. Lim, S. Ö. Arık, N. Loeff, and T. Pfister. Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37(4):1748–1764, 2021.

[32] J. M. Bates and C. W. J. Granger. The combination of forecasts. *Operational Research Quarterly*, 20(4):451–468, 1969.

[33] R. T. Clemen. Combining forecasts: A review and annotated bibliography. *International Journal of Forecasting*, 5(4):559–583, 1989.

[34] W. L. Hamilton, R. Ying, and J. Leskovec. Inductive representation learning on large graphs. *Advances in Neural Information Processing Systems 30*, 2017.

[35] A. Veit, S. Belongie, and T. Karaletsos. Conditional similarity networks. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 830–838, 2017.

[36] A. Das, W. Kong, R. Sen, and Y. Zhou. A decoder-only foundation model for time-series forecasting. *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.

[37] Google Research. TimesFM repository: TimesFM 2.5 (200M) and covariate support through XReg. https://github.com/google-research/timesfm, 2025; checkpoint `google/timesfm-2.5-200m-pytorch`, package `timesfm` 2.0.2.

[38] A. F. Ansari, L. Stella, C. Turkmen, X. Zhang, P. Mercado, H. Shen, O. Shchur, S. S. Rangapuram, S. Pineda Arango, S. Kapoor, J. Zschiegner, D. C. Maddix, H. Wang, M. W. Mahoney, K. Torkkola, A. G. Wilson, M. Bohlke-Schneider, and Y. Wang. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024.

[39] K. Rasul, A. Ashok, A. R. Williams, H. Ghonia, R. Bhagwatkar, A. Khorasani, M. J. Darvishi Bayazi, G. Adamopoulos, R. Riachi, N. Hassen, M. Biloš, S. Garg, A. Schneider, N. Chapados, A. Drouin, V. Zantedeschi, Y. Nevmyvaka, and I. Rish. Lag-Llama: Towards foundation models for probabilistic time series forecasting. arXiv:2310.08278, 2023.

[40] G. Woo, C. Liu, A. Kumar, C. Xiong, S. Savarese, and D. Sahoo. Unified training of universal time series forecasting transformers. *Proceedings of the 41st International Conference on Machine Learning*, PMLR 235:53140–53164, 2024.

[41] A. Garza, C. Challu, and M. Mergenthaler-Canseco. TimeGPT-1. arXiv:2310.03589, 2023.

[42] D. Ha and J. Schmidhuber. Recurrent world models facilitate policy evolution. *Advances in Neural Information Processing Systems 31*, 2018.

[43] D. Hafner, T. Lillicrap, I. Fischer, R. Villegas, D. Ha, H. Lee, and J. Davidson. Learning latent dynamics for planning from pixels. *Proceedings of the 36th International Conference on Machine Learning*, PMLR 97:2555–2565, 2019.

[44] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap. Mastering diverse domains through world models. arXiv:2301.04104, 2023.

[45] Y. LeCun. A path towards autonomous machine intelligence. OpenReview preprint, 2022.

[46] M. Assran, Q. Duval, I. Misra, P. Bojanowski, P. Vincent, M. Rabbat, Y. LeCun, and N. Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 15619–15629, 2023.

[47] M. Assran, A. Bardes, D. Fan, Q. Garrido, R. Howes, M. Komeili, M. Muckley, A. Rizvi, C. Roberts, K. Sinha, A. Zholus, S. Arnaud, A. Gejji, A. Martin, F. R. Hogan, D. Dugas, P. Bojanowski, V. Khalidov, P. Labatut, F. Massa, M. Szafraniec, K. Krishnakumar, Y. Li, X. Ma, S. Chandar, F. Meier, Y. LeCun, M. Rabbat, and N. Ballas. V-JEPA 2: Self-supervised video models enable understanding, prediction and planning. arXiv:2506.09985, 2025.

[48] J.-B. Grill, F. Strub, F. Altché, C. Tallec, P. H. Richemond, E. Buchatskaya, C. Doersch, B. Avila Pires, Z. D. Guo, M. Gheshlaghi Azar, B. Piot, K. Kavukcuoglu, R. Munos, and M. Valko. Bootstrap your own latent: A new approach to self-supervised learning. *Advances in Neural Information Processing Systems 33*, pages 21271–21284, 2020.

[49] A. Bardes, J. Ponce, and Y. LeCun. VICReg: Variance-invariance-covariance regularization for self-supervised learning. *International Conference on Learning Representations (ICLR)*, 2022.

[50] R. Balestriero and Y. LeCun. LeJEPA: Provable and scalable self-supervised learning without the heuristics. arXiv:2511.08544, 2025.

[51] L. Maes, Q. Le Lidec, D. Scieur, Y. LeCun, and R. Balestriero. LeWorldModel: Stable end-to-end joint-embedding predictive architecture from pixels. arXiv:2603.19312, 2026.

[52] N. Hansen, H. Su, and X. Wang. TD-MPC2: Scalable, robust world models for continuous control. *International Conference on Learning Representations (ICLR)*, 2024.

[53] Y.-G. Zhang, T. Du, Q. Zhang, and Y. Wang. DWM: Separating world effects from actions in latent world models. arXiv:2607.18715, 2026.

[54] J. Qiu, Z. Chen, H. Cao, J. Shi, J. Huo, and Y. Gao. AD-WM: Action-discriminative world models for counterfactual model predictive control. arXiv:2609.30264, 2026.

[55] L. Maes, Q. Le Lidec, L. Facury, N. Massaudi, A. Chaurasia, F. Capuano, R. Gao, T. Gillin, D. Haramati, D. Scieur, Y. LeCun, and R. Balestriero. stable-worldmodel: A platform for reproducible world modeling research and evaluation. arXiv:2605.21800, 2026.

[56] Y. Zhu, X. Qiu, H. Cheng, W. Shen, C. Guo, B. Yang, J. Hu, and C. S. Jensen. WorldTS: World modeling for multimodal covariate-aware time series forecasting. arXiv:2609.31162, 2026.

[57] K. Iyer. M-JEPA: joint-embedding predictive pretraining for molecular property prediction, code and evaluation protocol. https://github.com/Kar488/M-JEPA.

[58] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio. Learning phrase representations using RNN encoder–decoder for statistical machine translation. *Proceedings of EMNLP*, 2014.

[59] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. Attention is all you need. *Advances in Neural Information Processing Systems 30*, 2017.

[60] J. L. Ba, J. R. Kiros, and G. E. Hinton. Layer normalization. arXiv:1607.06450, 2016.

[61] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(56):1929–1958, 2014.

[62] I. Loshchilov and F. Hutter. Decoupled weight decay regularization. *International Conference on Learning Representations (ICLR)*, 2019.