Against the model itself
C1 holds by construction
C2 changes when weights are randomised
C3 same across seeds
C8 plan readable from the state
Against what sold
C5 predicted lift against sales in promoted and price-cut weeks
One forecast, explained
baseline (nothing planned)
× lift: promotion, price, display, ...
Against common sense
C4 a deeper cut never lowers the forecast
Question: do versions whose explanations pass more checks also forecast better?
Against another model and a cause
C6 agrees with LightGBM (on/off, TreeSHAP)
C7 matches the measured response to zone price deviations